

# Competition-Level Problems Are Effective Evaluators of LLMs

Yiming Huang<sup>1\*†</sup>, Zhenghao Lin<sup>2\*†</sup>, Xiao Liu<sup>1‡</sup>, Yeyun Gong<sup>1‡</sup>, Shuai Lu<sup>1</sup>,  
Fangyu Lei, Yaobo Liang<sup>1</sup>, Yelong Shen<sup>3</sup>, Chen Lin<sup>2‡</sup>, Nan Duan<sup>1</sup>, Weizhu Chen<sup>3</sup>

<sup>1</sup> Microsoft Research Asia, <sup>2</sup> Xiamen University, <sup>3</sup> Microsoft Azure AI  
{xiaoliu2,yegong}@microsoft.com, chenlin@xmu.edu.cn

## Abstract

Large language models (LLMs) have demonstrated impressive reasoning capabilities, yet there is ongoing debate about these abilities and the potential data contamination problem recently. This paper aims to evaluate the reasoning capacities of LLMs, specifically in solving recent competition-level programming problems in Codeforces, which are expert-crafted and unique, requiring deep understanding and robust reasoning skills. We first provide a comprehensive evaluation of GPT-4’s perceived zero-shot performance on this task, considering various aspects such as problems’ release time, difficulties, and types of errors encountered. Surprisingly, the perceived performance of GPT-4 has experienced a cliff like decline in problems after September 2021 consistently across all the difficulties and types of problems, which shows the potential data contamination, as well as the challenges for any existing LLM to solve unseen complex reasoning problems. We further explore various approaches such as fine-tuning, Chain-of-Thought prompting and problem description simplification, unfortunately none of them is able to consistently mitigate the challenges. Through our work, we emphasize the importance of this excellent data source for assessing the genuine reasoning capabilities of LLMs, and foster the development of LLMs with stronger reasoning abilities and better generalization in the future.

## 1 Introduction

The rise of LLMs has generated significant interest in the artificial intelligence community. These models, notably GPT-4 (OpenAI, 2023), have displayed impressive reasoning capabilities that are being harnessed in various fields (Bubeck et al., 2023). However, questions<sup>1</sup> have been raised about

\* Equal contribution.

† This work was done during their internship at MSRA.

‡ Corresponding authors.

<sup>1</sup><https://twitter.com/keirp1/status/1724518513874739618>




Figure 1: The perceived zero-shot performance of GPT-4 sees a sharp decline on problems of varying difficulties (D1, D2 and D3 means easy, medium and difficult, respectively) in Codeforces after September 2021.

how to accurately evaluate the reasoning abilities of LLMs and the extent of data contamination issues (Mialon et al., 2023; Zhou et al., 2023).

Regarding these issues, our study aims to assess the reasoning capabilities of LLMs through their ability to generate algorithms for solving competition-level programming problems. **These questions are meticulously crafted by experts to form rigorous competitions. They possess high quality, are unique, and exhibit excellent discriminative ability. The testing cases are also meticulously prepared.** This necessitates that LLMs deduce the solution from the presented scenario, which requires a thorough understanding of algorithms, combined reasoning and coding skills, and strong problem-solving abilities. These problems thus present a significant challenge to both human coders and LLMs. Consequently, competition-level programming problems serve as effective tools for evaluating the two issues previously discussed: they assess the reasoning abilities of LLMs and, due to the strict problem selection process in competitions, reduce the likelihood of

data contamination in new problems.

Our research provides an in-depth analysis of the zero-shot performances of GPT-4 and other code LLMs on competition-level programming problems in Codeforces, considering factors such as the release time, problem difficulty, and the types of errors encountered. The main insights of our study include: (1) GPT-4 performs significantly worse on programming problems released after September 2021, casting doubt on its actual reasoning abilities. (2) GPT-4 shows limited capability to solve difficult problems, indicating potential weaknesses in complex problem-solving. (3) GPT-4 struggles with the first test case, suggesting errors may stem from its understanding of the problem at hand. (4) The related phenomenon can be also observed in other LLMs, indicating that insufficient reasoning ability may be a common problem.

To explore possible ways to enhance the zero-shot performances of these LLMs on competition-level programming problems, we investigate several methods to improve performance on unseen problems. These include supervised fine-tuning with CodeLlama (Rozière et al., 2023) and DeepSeek-Coder (AI, 2023), Chain-of-Thought prompting (Wei et al., 2022), and problem statement simplification. However, none of these methods consistently mitigate the issue or result in noticeable performance improvements, particularly for more difficult problems. This finding indicates that the difficult and unseen programming problems are effective evaluators of LLMs.

Overall, the primary contributions of this study lie in proposing and validating that recent competition-level programming problems serve as an excellent data source for assessing the genuine reasoning capabilities of LLMs. We aim to foster further research in this field by innovating new approaches to address the challenge of complex reasoning problems in LLMs and by establishing reliable evaluation benchmarks for AI models that minimize the risk of data contamination.

## 2 Problem Setup

### 2.1 Competition-level Programming

Competition-level programming presents a unique arena for testing and developing the reasoning abilities of AI models. In these competitions, participants must design algorithms and implement them in programming languages like C++ and Java. Accepted programs must satisfy stringent testing con-

ditions, including producing outputs that exactly match with test cases, executing within memory limits, and terminating within time constraints. In contrast to prior works (Chen et al., 2021a; Austin et al., 2021; Cassano et al., 2023) focusing on basic coding abilities, competition-level programming problems require advanced reasoning and mathematical modeling skills, essential for AI.

Unlike the previous works that focused on LeetCode<sup>2</sup> (Bubeck et al., 2023; Shen et al., 2023; Sakib et al., 2023), we follow AlphaCode (Li et al., 2022) and choose Codeforces<sup>3</sup>. Codeforces is universally acknowledged by competitors and enthusiasts in the International Collegiate Programming Competition<sup>4</sup> (ICPC) and the International Olympiad in Informatics<sup>5</sup> (IOI) as a popular and suitable platform for developing abilities for algorithm contests. The regular contests hosted on this platform are crafted by human experts, and contain plenty of intricate programming problems and contests of high quality. These contests come with comprehensive and robust test cases and exhibit a low degree of problem overlap. The unique nature of these contest problems makes it highly unlikely to find similar content on the internet before the competition concludes. As a result, utilizing specific time-segmented datasets, like those from contests conducted post the introduction of LLMs, serves as an effective strategy to prevent data contamination (Zhou et al., 2023).

Codeforces employs the Elo rating system<sup>6</sup> to rank its users and problems, categorizing all problems into 28 distinct difficulties, ranging from 800 to 3500. Compared to commonly utilized metrics such as the ratio of accepted submissions or users, this difficulty rating mechanism is more suitable as it is based on the ranking and performance of the participants during the competition.<sup>7</sup> Hence, it is not subject to inaccuracies stemming from temporal changes, repeated submissions, plagiarism, and other potential distortions.

### 2.2 Problem Definition

Figure 2 presents an example of the problem statement  $\pi$ . The input of LLM is instantiated with the problem statement  $\pi$  and a prompt  $\rho$  (like  $\rho_1$  in

<sup>2</sup><https://leetcode.com/>

<sup>3</sup><https://codeforces.com/>

<sup>4</sup><https://icpc.global/>

<sup>5</sup><https://ioinformatics.org/>

<sup>6</sup><https://codeforces.com/blog/entry/102>

<sup>7</sup><https://codeforces.com/blog/entry/62865>


<p><b>Kefa and Park</b></p> <p>The semester is already ending, so Danil made an effort and decided to visit a lesson on harmony analysis to know how does the professor look like, at least. Danil was very bored on this lesson until the teacher gave the group a simple task: find 4 vectors in 4-dimensional space, such that every coordinate of every vector is 1 or -1 and any two vectors are orthogonal. Just as a reminder, two vectors in n-dimensional space are considered to be orthogonal if and only if their scalar product is equal to zero, that is: <math>\sum_{i=1}^n v_i \cdot w_i = 0</math>. Danil quickly managed to come up with the solution for this problem and the teacher noticed that the problem can be solved in a more general case for 2k vectors in 2k-dimensional space. When Danil came home, he quickly came up with the solution for this problem. Can you cope with it?</p> <p><b>Input</b></p> <p>The only line of the input contains a single integer k (<math>0 \leq k \leq 9</math>).</p>	<p><b>Output</b></p> <p>Print 2k lines consisting of 2k characters each. The j-th character of the i-th line must be equal to '*' if the j-th coordinate of the i-th vector is equal to -1, and must be equal to '.' if it's equal to +1. It's guaranteed that the answer always exists. If there are many correct answers, print any.</p> <p><b>Time limit per test</b></p> <p>3 seconds</p> <p><b>Memory limit per test</b></p> <p>256 megabytes</p> <p><b>Example</b></p> <p><b>input</b></p> <p>2</p> <p><b>output</b></p> <p>++**\n+*+*\n++++\n+**+\n</p> <p><b>Note</b></p> <p>none</p>	<pre>// generated by GPT4 #include&lt;bits/stdc++.h&gt; using namespace std; vector&lt;vector&lt;int&gt;&gt;&gt; v(1, vector&lt;int&gt;(1, 1)); void solve(int k) {     if(k == 0) return;     int sz = 1 &lt;&lt; (k - 1);     for(int i = 0; i &lt; sz; i++) {         for(int j = 0; j &lt; sz; j++) {             v[i].push_back(v[i][j]);             v[i].push_back(v[i][j]);         }     }     for(int i = 0; i &lt; sz; i++) {         for(int j = 0; j &lt; sz; j++) {             v[i + sz][j] = -1;         }     }     solve(k - 1); } int main() {     int k;     cin &gt;&gt; k;     solve(k);     int sz = 1 &lt;&lt; k;     for(int i = 0; i &lt; sz; i++) {         for(int j = 0; j &lt; sz; j++) {             cout &lt;&lt; (v[i][j] == 1 ? '+' : '*');         }         cout &lt;&lt; "\n";     }     return 0; }</pre> <div style="display: flex; justify-content: space-between; align-items: center;">  <span>Online Judge</span> <span>AC/WA/CE/...</span> </div>
---	---	---

Figure 2: The figure depicts the problem statement (left), comprising a problem set in a narrative context for participants to decipher, detailed input and output format specifications, and one or more example input-output pairs. In some cases, additional notes may be provided to assist competitors in understanding these example tests. This information is fed into the LLM, aiming to generate relevant code (right). The generated code is then submitted to an online judge for correctness evaluation.

Table 7). The LLM  $\Gamma$  takes the input to generate the code as  $\alpha = \Gamma(\rho(\pi))$ . The generated code  $\alpha$  is then evaluated by an online judge (OJ). The evaluation process can be summarized in the following equation:

$$OJ(\alpha) = OJ(\Gamma(\rho(\pi))) \in \{AC, WA, CE, \dots\}$$

In this equation,  $\Gamma(\rho(\pi))$  denotes the code generated by LLM with the prompt  $\rho$ . The OJ platform then rigorously assesses the code for its correctness, computational efficiency, and adherence to specified input/output formats. With an extensive testing mechanism, the platform employs a wide range of test cases and hidden scenarios to ensure the code's robustness across diverse scenarios. The platform provides a spectrum of outcomes,  $OJ(\Gamma(\rho(\pi)))$ , offering a holistic evaluation of the code's performance. This includes results such as Accepted (AC), Wrong Answer (WA), and Compilation Error (CE), among others.

### 2.3 Dataset Collection

The dataset is compiled from the Codeforces website, extracting all publicly available problem statements from completed contests spanning February 2010 through November 2023. For simplicity, problems requiring interaction, featuring non-standard

Metric	Definition
ACC#G	Proportion of accepted solutions using greedy sampling (temperature $t = 0$ ).
ACC#GN	The number of accepted solutions using greedy sampling (temperature $t = 0$ ) within the sliding window.
ACCk#n	Proportion of problems with $k$ or more accepted solution with top-p samplings ( $t = 0.7, p = 0.95$ ) for $n$ times.
pass@k	Estimated proportion of problems with at least one accepted solution.

Table 1: Definitions of evaluation metrics.

input/output formats, or incompatible with C++ submission are excluded.

The analysis is confined to problems with difficulty levels ranging from 800 to 2400. Based on their difficulty levels, the dataset is divided into three subsets: **D1** (800-1100 difficulty, 1683 problems), **D2** (1200-1600 difficulty, 1821 problems), and **D3** (1700-2400 difficulty, 1453 problems). These problems encompass more than 20 distinct categories of algorithms, as illustrated in Table 5. This diversity in problem types further enhances the comprehensiveness of the dataset and enables a comprehensive assessment of GPT-4's problem-solving abilities across a wide range of competition-level programming problems.

Metric	D1			D2			D3		
	Time1	Time2	$\Delta$	Time1	Time2	$\Delta$	Time1	Time2	$\Delta$
ACC#G	81.42%	11.73%	-69.69%	43.72%	0.00%	-43.72%	11.41%	0.00%	-11.41%
pass@1	78.11%	10.54%	-67.57%	42.38%	0.61%	-41.77%	9.45%	0.18%	-9.27%
ACC1#1	78.05%	9.38%	-68.68%	43.37%	0.00%	-43.37%	8.48%	0.00%	-8.48%
ACC1#5	94.03%	20.09%	-73.94%	69.02%	3.06%	-65.96%	21.24%	0.88%	-20.36%
ACC2#5	88.34%	11.83%	-76.51%	54.41%	0.00%	-54.41%	12.36%	0.00%	-12.36%
ACC3#5	81.82%	9.38%	-72.44%	42.42%	0.00%	-42.42%	7.51%	0.00%	-7.51%

Table 2: Performance of GPT-4 on different groups of problems: Time1 is the problems released from October 2010 to September 2021, and Time2 is the problems released from October 2021 to November 2023.

## 2.4 Experiment Details

In Codeforces, each problem belongs to a contest. Once the contest concludes, the problems are disclosed and become publicly submittable. Therefore, we submit the solutions to the contests that have concluded for evaluation.

To evaluate the results, we employ ACC#G, ACC#GN, ACC $k$ # $n$  and pass@ $k$  as defined in Table 1. Specifically, for ACC $k$ # $n$  metric, we consider two settings: (1)  $k = n = 1$  and (2)  $k \in \{1, 2, 3\}$  with  $n = 5$ . Following Codex (Chen et al., 2021b), pass@ $k$  is computed as

$$\text{pass@}k := \mathbb{E}_{\text{Problems}} \left[ 1 - \frac{\binom{n-c}{k}}{\binom{n}{k}} \right]$$

where  $n$  is defined as the total number of generated samples per problem used for evaluation, and  $c$  represents the count of correct samples out of  $n$  that have successfully passed the unit tests. Here we use  $k = 1$  and  $n = 5$  for pass@ $k$ .

In our experiment, we follow the zero-shot setting. To select an appropriate prompt, we conduct preliminary experiments with three prompts,  $\rho_1$ ,  $\rho_2$ , and  $\rho_3$ , as listed in Table 7, using two subsets of **D1** problems: one from February to December 2010 and the other from January to October 2023, each comprising approximately 100 problems. The standard deviations are 0.015 and 0.018, respectively, indicating consistent performance. Therefore, we choose  $\rho_1$  as the prompt in the subsequent experiments. Furthermore, we employ a sliding window approach for all temporal analyses to smooth the data, addressing the sporadic release schedule of the problems. This ensures a sufficient number of test problems at each time point, using a window size of 51 (25 before and 25 after the time point).

## 3 Insights and Implications

### 3.1 Faltering on Unseen Problems

In this section, we delve into a temporal analysis of GPT-4’s performance on programming prob-




Figure 3: Random sampling enhances the probability of generating correct solutions on previously encountered problems, but offers no assistance for unseen problems.

lems. Figure 1 illustrates GPT-4’s performance using the ACC#G metric. On problems released prior to September 2021, GPT-4 exhibits minor fluctuations at different levels across problems of varying difficulty. However, for problems released after September 2021, a significant deviation from the normal fluctuation range is observed. Interestingly, this timing coincides with the cut-off date for the GPT-4 training data as announced by OpenAI<sup>8</sup>. We then calculate the average performance on problems before and after September 2021, as shown in Table 2. On **D1** problems, GPT-4’s ACC#G plummets from 81.42% to 11.73%, marking a stark decrease of 69.69%. Even more strikingly, the ACC#G drops to 0.00% on both **D2** and **D3** problems, from 43.72% and 11.41%, respectively. To validate the reliability of the conclusion, we also calculate the pass@1 metric, which exhibits a similar trend. This observation raises thought-provoking questions about the severity of the drop and the correlation between the data cut-off date and the performance decline.

<sup>8</sup><https://platform.openai.com/docs/models/gpt-4-and-gpt-4-turbo>



To explore the model’s potential to generate correct solutions, we perform random sampling multiple times and calculate the pass rate. The average pass rate are shown in Table 2. As observed, multiple samplings can enhance the chances of generating a correct solution. For instance, on the unseen simple **D1** problems, ACC1#5 improved by 10.71% compared to ACC1#1. However, across all problems, the performance gap before and after the cut-off date is more pronounced for ACC1#5 than for both ACC1#1 and ACC#G. Figure 3 depicts the performance on **D2** problems over time. A notable decline in performance metrics is observed around September 2021. This observation underscores the challenges that LLMs, including the advanced GPT-4, face in addressing unseen programming problems without similar pretraining data.

The observed decline in performance on problems outside the model’s training range may stem from limitations in reasoning and generalization. As highlighted by [Yadlowsky et al. \(2023\)](#), when confronted with problems beyond their pretraining data, transformer models exhibit various failure modes and their generalization abilities deteriorate, even for simple problems. Similarly, [Lu et al. \(2023\)](#) suggest that the exceptional abilities of large language models primarily stem from in-context learning, and do not necessarily reflect the emergence of reasoning abilities.

The observed performance drop on unseen problems raises serious questions about GPT-4’s intrinsic reasoning and generalization capabilities. This suggests a potential over-reliance on pattern recognition and reproduction from training, as opposed to grasping underlying principles and applying them to novel problems. This observation aligns with recent debates on large models’ data memorization tendencies ([Carlini et al., 2023](#); [Yang et al., 2023](#)). Therefore, future evaluations should prioritize the minimization of overlap between testing and training data to accurately assess a model’s reasoning abilities, rather than simply its capacity for memorization. Furthermore, it’s crucial to explore methods that enhance model generalization and reduce reliance on pre-training data.

### 3.2 Limited Ability to Solve Difficult Problems

This section provides an analysis of performance in relation to the problem difficulty. The results of ACC#G for problems with different difficulties are reported for two distinct periods: from October




Figure 4: For problems released before September 2021, GPT-4’s ACC#G showed a negative linear correlation with difficulty, followed by consistently poor performance afterwards.

2010 to September 2021, and from October 2021 to November 2023, as illustrated in Figure 4.

For the results from October 2010 to September 2021, we calculate Pearson correlation coefficient ( $r = -0.97$ ) and the Kendall rank correlation coefficient ( $\tau = -0.88$ ), which indicate strong linear correlations. Notably, when the difficulty level reaches 2400 (indicating greater challenge than approximately 57% of the problems on Codeforces), the ACC#G drops to zero. However, from October 2021 to November 2023, ACC#G shows a dramatic decrease across all difficulty levels.

These findings reveal a significant limitation in the ability of GPT-4 to handle extremely complex problems. Despite its vast knowledge on code and algorithms, GPT-4 lacks of the competence in solving very challenging problems, particularly those with higher difficulty levels, even in the context of previously encountered problems. This indicates a potential area for further improvement and development in future iterations of the model.

### 3.3 Struggling with The First Test Case

In this section, we gather and analyze the errors returned by GPT-4 upon submission to the Codeforces website, as outlined in Table 6. The most common error is "Wrong answer on test 1", which on average accounts for 70% of the observed errors. Test 1 is the first test case, which almostly corresponds to or properly includes the example test case provided in the problem statement. This suggests that the model often struggles at the very beginning of problem-solving, possibly due to difficulties in




Figure 5: Error categories in GPT-4’s solutions on problems released from 2010 to 2023.




Figure 6: ACC#GN of CodeLlama and DeepSeek-Coder on **D1** problems.

understanding the problem’s requirements or generating a correct solution based on the given test case. As depicted in Figure 5, there is a significant increase in the proportion of "Wrong answer on test 1" errors for problems released between 2021 and 2023. This suggests that GPT-4 is more likely to face challenges in understanding and reasoning during at the onset of tackling unseen problems.

Other types of errors account for a smaller proportion, with an average of 10%. They have shown little variation over time. This indicates that GPT-4 demonstrates strong fundamental code-writing capabilities of generating high-quality code.

### 3.4 Similar Phenomenons of Other Code LLMs

We investigate whether the perceived performance degradation on unseen programming problems is

Problem Release Date	CodeLlama	DeepSeek-Coder
Before 2023.3	10.30%	32.74%
After 2023.3	4.52% (-5.78%)	9.03% (-23.71%)

Table 3: Comparison of ACC#G between CodeLlama and DeepSeek-Coder on **D1** problems before and after March 2023.

Model	2020.1-2021.9	2021.9-2023.10
GPT-4	73.19% (+50.52%)	11.53% (-0.97%)
DeepSeek-Coder	22.67%	12.50%

Table 4: Comparison of ACC#G between GPT-4 and DeepSeek-Coder over time intervals, on **D1** problems.

observed for other popular code LLMs, such as CodeLlama-34B-Instruct (Rozière et al., 2023) and DeepSeek-Coder-33B-Instruct (AI, 2023).

We conduct tests on CodeLlama and DeepSeek-Coder using **D1** problems, following the settings in §2.3, and the results are shown in Figure 6. The experimental results indicate that CodeLlama consistently underperforms compared to DeepSeek-Coder on **D1** problems. Furthermore, the performance of DeepSeek-Coder on **D1** problems has been declining with the progression of the problem release date. The ACC#GN of DeepSeek-Coder has declined to a level that is on par with CodeLlama when dealing with newly released problems, as highlighted in the red area of Figure 6.

To precisely and intuitively detect this phenomenon, we calculate the ACC#G of CodeLlama and DeepSeek-Coder on **D1** problems, both before and after March 2023, and present the results in Table 3. The results reveal a significant difference in the average accuracy of CodeLlama and DeepSeek-Coder before and after March 2023. Regarding the magnitude of the decrease, DeepSeek-Coder, which previously exhibited superior performance, demonstrates a more pronounced decline, with acceptance rates falling below 10% after March 2023. Considering the release dates of CodeLlama and DeepSeek-Coder, we speculate that most of the programming problems after March 2023 are novel to them, which suggests that they also not be able to perform well on unseen programming problems like GPT4 does. This finding indicates that a fundamental limitation of current code LLMs in generalizing effectively to complex reasoning tasks.

### 3.5 Evaluation Hallucination of LLMs

To further analyze the phenomenon, we compare GPT-4 with DeepSeek-Coder on **D1** problems as shown in Figure 7 and Table 4.




Figure 7: Comparison of ACC#GN of GPT-4 and DeepSeek-Coder on **D1** problems after 2020.

It is noteworthy that while GPT-4 surpasses DeepSeek-Coder in terms of performance on problems that were released prior to September 2021, an unexpected observation is that DeepSeek-Coder exhibits a performance that is on par with GPT-4 when it comes to tackling problems that were released after September 2021. Considering the previous work (Yang et al., 2023; Zhou et al., 2023), although GPT-4 may perform particularly well on some previously seen problems due to its powerful capacity, it cannot be well generalized on unseen programming problems, and its performance is not significantly different from DeepSeek-Coder, which is specifically trained for code. This phenomenon merits attention, which is termed as “evaluation hallucination”.

Hence, a more equitable evaluation strategy would be to select evaluation sets that all the models have not previously encountered. However, finding such data adhering to stringent conditions is challenging, as LLMs are typically pre-trained on extensive corpora containing diverse content, leading to the potential issue of data contamination. Therefore, if we could devote more attention to the data source and timeline of the evaluation sets, such as the problems in Codeforces, it could potentially mitigate the effects of evaluation hallucination.

## 4 One Step Forward

In this section, we explore some approaches to mitigate the poor performance on unseen problems.

### 4.1 Finetuning

Fine-tuning is a commonly used method to improve the performance on a specific downstream task. Specifically, we use Description2Code (OpenAI and Sutskever, 2016) as fine-tuning dataset, which




Figure 8: Comparison of ACC#GN on **D1** problems before and after fine-tuning CodeLlama.




Figure 9: Comparison of ACC#GN on **D1** problems before and after fine-tuning DeepSeek-Coder.

contains approximately 7,000 problems released before 2017. We use approximately 10 solutions for each problem, resulting 70,000 pairs of input-output sequences, and employ them for fine-tuning both CodeLlama and DeepSeek-Coder in a supervised manner.

As shown in Figure 8 and Figure 9, we compare the performances of the models before and after fine-tuning on **D1** problems. We observe that, even after fine-tuning with the same type of data, CodeLlama and DeepSeek-Coder do not exhibit improved performance on recent problems, particularly those post-2022. The significant improvement in ACC#GN before 2017 may result from the models recalling relevant or identical programming problems, rather than mastering the underlying reasoning logic, leading to their inability to adapt well to new programming challenges. Therefore, simple fine-tuning does not effectively enhance the models’ performance on new programming problems.

### 4.2 Chain-of-Thought Prompting

In this section, we explore the application of Chain-of-Thought (CoT) prompting (Wei et al., 2022) to

competition-level programming problems. CoT involves prompting GPT-4 to generate an explanation of the algorithm before coding, denoted as  $\rho_{cot}$  in Table 7. We conduct experiments on both the **D1** and **D3** problems released after October 2021. For **D1** problems, employing CoT increases the ACC#G from 11.54% to 16.21%, demonstrating a noticeable improvement. However, for **D3** problems, using CoT fails to yield any improvement, leaving the ACC#G at 0.00%. This suggests that while CoT facilitates some improvement for simple **D1** problems, it is ineffective for the complex reasoning challenges presented by **D3** problems.

### 4.3 Problem Statement Simplification

Intuitively, even experienced programming competition competitors require time to understand problem statements. Therefore, we conduct a simple experiment to assess whether comprehension of problem statements hinders LLMs’ ability to excel at programming problems. We first instruct GPT-4 to simplify the problem statement with  $\rho_{sip}$  and then generate the code with  $\rho_{sipgen}$  as shown in Table 7. The results are also evaluated on both the **D1** and **D3** problems released after October 2021. However, for **D1** problems, using the simplified problem statement even brings a slight decline in ACC#G from 11.54% to 11.14%. And the ACC#G for **D3** problems still remains at 0.00%. Consequently, the challenge of genuinely improving the model’s reasoning ability and enhancing its performance on unseen problems represents a significant direction for future research.

## 5 Related Work

**Code LLMs.** Code intelligence is an important topic in AI research. Recently, code LLMs (Zhang et al., 2023b) have received widespread attention. Commercial LLMs (OpenAI, 2023) have achieved tremendous success. Meanwhile, research on open-source code LLMs is also thriving, such as CodeLlama (Rozière et al., 2023), StarCoder (Li et al., 2023), CodeGeeX (Zheng et al., 2023), CodeFuse (Di et al., 2023), WizardCoder (Luo et al., 2023) and Lemur (Xu et al., 2023).

**Reasoning on Code.** Programming competition is a specialized domain within the broader landscape of programming problems. Unlike simpler tasks on code, such as HumanEval (Chen et al., 2021a), MBPP (Austin et al., 2021), MultiPLE (Cassano et al., 2023), competition-level pro-

gramming problems necessitate an advanced understanding of data structures, algorithms, and problem-solving techniques. Enabling models to solve human-designed algorithmic competition problems represents a meaningful research direction, as it reflects the models’ integrated capabilities in reasoning, coding, and problem-solving. AlphaCode (Li et al., 2022) simulate evaluations on 10 programming competitions on the Codeforces platform, which is the first work in this topic. ALGO (Zhang et al., 2023a) can integrate with any existing code LLMs in a model-agnostic manner, enhancing its code generation performance.

**Reasoning on Other Subjects.** Researchers have proposed many benchmarks requiring various reasoning skills, including commonsense reasoning (Talmor et al., 2018; Geva et al., 2021), numerical reasoning (Dua et al., 2019), multi-hop reasoning (Yang et al., 2018), arithmetic reasoning (Patel et al., 2021; Cobbe et al., 2021), structured reasoning (Yu et al., 2018; Lei et al., 2023), inductive reasoning (Sinha et al., 2019) and logical reasoning (Yu et al., 2020). LLMs are also widely used in scientific research in other fields (Wang et al., 2023), such as physics (Yeadon and Halliday, 2023), chemistry (Castro Nascimento and Pimentel, 2023; Bran et al., 2023), etc.

## 6 Conclusion

In this study, we utilize competition-level programming problems from Codeforces to analyze the reasoning capabilities of LLMs. We find a significant decrease in perceived performance of GPT-4 on unseen problems, consistent across a range of difficulties, problem types, and experimental settings. This decrease highlights concerns of data contamination in benchmarks and the need for unseen tasks to properly assess LLMs’ reasoning ability with complex challenges. Our research also extends these insights to other open-source LLMs, revealing the common difficulties these models face with complex, previously unencountered reasoning tasks. This is indicative of the LLMs’ intrinsic limitations in reasoning. As a primary probe, we explore several straightforward strategies, but none of them consistently mitigated the issues. Through our work, we hope to emphasize the critical need for robust datasets to accurately evaluate LLMs’ reasoning abilities and to inspire advancements in LLMs that demonstrate improved reasoning abilities.



## References

- DeepSeek AI. 2023. Deepseek coder: Let the code write itself. <https://github.com/deepseek-ai/DeepSeek-Coder>.
- Jacob Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, Ellen Jiang, Carrie Cai, Michael Terry, Quoc Le, et al. 2021. Program synthesis with large language models. *arXiv preprint arXiv:2108.07732*.
- Andres M Bran, Sam Cox, Andrew D White, and Philippe Schwaller. 2023. Chemcrow: Augmenting large-language models with chemistry tools. *arXiv preprint arXiv:2304.05376*.
- Sébastien Bubeck, Varun Chandrasekaran, Ronen Eldan, Johannes Gehrke, Eric Horvitz, Ece Kamar, Peter Lee, Yin Tat Lee, Yuanzhi Li, Scott Lundberg, et al. 2023. Sparks of artificial general intelligence: Early experiments with gpt-4. *arXiv preprint arXiv:2303.12712*.
- Nicholas Carlini, Daphne Ippolito, Matthew Jagielski, Katherine Lee, Florian Tramèr, and Chiyuan Zhang. 2023. Quantifying memorization across neural language models. In *The Eleventh International Conference on Learning Representations*.
- Federico Cassano, John Gouwar, Daniel Nguyen, Sydney Nguyen, Luna Phipps-Costin, Donald Pinckney, Ming-Ho Yee, Yangtian Zi, Carolyn Jane Anderson, Molly Q Feldman, et al. 2023. Multipl-e: a scalable and polyglot approach to benchmarking neural code generation. *IEEE Transactions on Software Engineering*.
- Cayque Monteiro Castro Nascimento and André Silva Pimentel. 2023. Do large language models understand chemistry? a conversation with chatgpt. *Journal of Chemical Information and Modeling*, 63(6):1649–1655.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. 2021a. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. 2021b. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, et al. 2021. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*.
- Peng Di, Jianguo Li, Hang Yu, Wei Jiang, Wenting Cai, Yang Cao, Chaoyu Chen, Dajun Chen, Hongwei Chen, Liang Chen, Gang Fan, Jie Gong, Zi Gong, Wen Hu, Tingting Guo, Zhichao Lei, Ting Li, Zheng Li, Ming Liang, Cong Liao, Bingchang Liu, Jiachen Liu, Zhiwei Liu, Shaojun Lu, Min Shen, Guangpei Wang, Huan Wang, Zhi Wang, Zhaogui Xu, Jiawei Yang, Qing Ye, Gehao Zhang, Yu Zhang, Zelin Zhao, Xunjin Zheng, Hailian Zhou, Lifu Zhu, and Xianying Zhu. 2023. *Codefuse-13b: A pretrained multi-lingual code large language model*. *CoRR*, abs/2310.06266.
- Dheeru Dua, Yizhong Wang, Pradeep Dasigi, Gabriel Stanovsky, Sameer Singh, and Matt Gardner. 2019. Drop: A reading comprehension benchmark requiring discrete reasoning over paragraphs. *arXiv preprint arXiv:1903.00161*.
- Mor Geva, Daniel Khashabi, Elad Segal, Tushar Khot, Dan Roth, and Jonathan Berant. 2021. Did aristotle use a laptop? a question answering benchmark with implicit reasoning strategies. *Transactions of the Association for Computational Linguistics*, 9:346–361.
- Fangyu Lei, Qian Liu, Yiming Huang, Shizhu He, Jun Zhao, and Kang Liu. 2023. S3eval: A synthetic, scalable, systematic evaluation suite for large language models. *arXiv preprint arXiv:2310.15147*.
- Raymond Li, Loubna Ben Allal, Yangtian Zi, Niklas Muennighoff, Denis Kocetkov, Chenghao Mou, Marc Marone, Christopher Akiki, Jia Li, Jenny Chim, Qian Liu, Evgenii Zheltonozhskii, Terry Yue Zhuo, Thomas Wang, Olivier Dehaene, Mishig Davaadorj, Joel Lamy-Poirier, João Monteiro, Oleh Shliazhko, Nicolas Gontier, Nicholas Meade, Armel Zebaze, Ming-Ho Yee, Logesh Kumar Umapathi, Jian Zhu, Benjamin Lipkin, Muhtasham Oblokulov, Zhiruo Wang, Rudra Murthy V, Jason Stiller, Siva Sankalp Patel, Dmitry Abulkhanov, Marco Zocca, Manan Dey, Zhihan Zhang, Nour Moustafa-Fahmy, Urvashi Bhattacharyya, Wenhao Yu, Swayam Singh, Sasha Luccioni, Paulo Villegas, Maxim Kunakov, Fedor Zhdanov, Manuel Romero, Tony Lee, Nadav Timor, Jennifer Ding, Claire Schlesinger, Hailley Schoelkopf, Jan Ebert, Tri Dao, Mayank Mishra, Alex Gu, Jennifer Robinson, Carolyn Jane Anderson, Brendan Dolan-Gavitt, Danish Contractor, Siva Reddy, Daniel Fried, Dzmitry Bahdanau, Yacine Jernite, Carlos Muñoz Ferrandis, Sean Hughes, Thomas Wolf, Arjun Guha, Leandro von Werra, and Harm de Vries. 2023. *StarCoder: may the source be with you!* *CoRR*, abs/2305.06161.
- Yujia Li, David Choi, Junyoung Chung, Nate Kushman, Julian Schrittwieser, Rémi Leblond, Tom Eccles, James Keeling, Felix Gimeno, Agustin Dal Lago, et al. 2022. Competition-level code generation with alphacode. *Science*, 378(6624):1092–1097.
- Sheng Lu, Irina Bigoulaeva, Rachneet Sachdeva, Harish Tayyar Madabushi, and Iryna Gurevych. 2023. Are emergent abilities in large language models just in-context learning? *arXiv preprint arXiv:2309.01809*.

- Ziyang Luo, Can Xu, Pu Zhao, Qingfeng Sun, Xubo Geng, Wenxiang Hu, Chongyang Tao, Jing Ma, Qingwei Lin, and Daxin Jiang. 2023. [Wizardcoder: Empowering code large language models with evol-instruct](#). *CoRR*, abs/2306.08568.
- Grégoire Mialon, Clémentine Fourrier, Craig Swift, Thomas Wolf, Yann LeCun, and Thomas Scialom. 2023. Gaia: a benchmark for general ai assistants. *arXiv preprint arXiv:2311.12983*.
- OpenAI. 2023. [Gpt-4 technical report](#).
- E. Caballero OpenAI and I. Sutskever. 2016. [Description2Code Dataset](#).
- Arkil Patel, Satwik Bhattamishra, and Navin Goyal. 2021. Are nlp models really able to solve simple math word problems? *arXiv preprint arXiv:2103.07191*.
- Baptiste Rozière, Jonas Gehring, Fabian Gloeckle, Sten Sootla, Itai Gat, Xiaoqing Ellen Tan, Yossi Adi, Jingyu Liu, Tal Remez, Jérémy Rapin, Artyom Kozhevnikov, Ivan Evtimov, Joanna Bitton, Manish Bhatt, Cristian Canton-Ferrer, Aaron Grattafiori, Wenhan Xiong, Alexandre Défossez, Jade Copet, Faisal Azhar, Hugo Touvron, Louis Martin, Nicolas Usunier, Thomas Scialom, and Gabriel Synnaeve. 2023. [Code llama: Open foundation models for code](#). *CoRR*, abs/2308.12950.
- Fardin Ahsan Sakib, Saadat Hasan Khan, and AHM Karim. 2023. Extending the frontier of chatgpt: Code generation and debugging. *arXiv preprint arXiv:2307.08260*.
- Bo Shen, Jiaxin Zhang, Taihong Chen, Daoguang Zan, Bing Geng, An Fu, Muhan Zeng, Ailun Yu, Jichuan Ji, Jingyang Zhao, et al. 2023. Pangu-coder2: Boosting large language models for code with ranking feedback. *arXiv preprint arXiv:2307.14936*.
- Koustuv Sinha, Shagun Sodhani, Jin Dong, Joelle Pineau, and William L Hamilton. 2019. Clutrr: A diagnostic benchmark for inductive reasoning from text. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 4506–4515.
- Alon Talmor, Jonathan Herzig, Nicholas Lourie, and Jonathan Berant. 2018. Commonsenseqa: A question answering challenge targeting commonsense knowledge. *arXiv preprint arXiv:1811.00937*.
- Hanchen Wang, Tianfan Fu, Yuanqi Du, Wenhao Gao, Kexin Huang, Ziming Liu, Payal Chandak, Shengchao Liu, Peter Van Katwyk, Andreea Deac, et al. 2023. Scientific discovery in the age of artificial intelligence. *Nature*, 620(7972):47–60.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, et al. 2022. Chain-of-thought prompting elicits reasoning in large language models. *Advances in Neural Information Processing Systems*, 35:24824–24837.
- Yiheng Xu, Hongjin Su, Chen Xing, Boyu Mi, Qian Liu, Weijia Shi, Binyuan Hui, Fan Zhou, Yitao Liu, Tianbao Xie, et al. 2023. Lemur: Harmonizing natural language and code for language agents. *arXiv preprint arXiv:2310.06830*.
- Steve Yadlowsky, Lyric Doshi, and Nilesch Tripathaneni. 2023. Pretraining data mixtures enable narrow model selection capabilities in transformer models. *arXiv preprint arXiv:2311.00871*.
- Shuo Yang, Wei-Lin Chiang, Lianmin Zheng, Joseph E Gonzalez, and Ion Stoica. 2023. Rethinking benchmark and contamination for language models with rephrased samples. *arXiv preprint arXiv:2311.04850*.
- Zhilin Yang, Peng Qi, Saizheng Zhang, Yoshua Bengio, William Cohen, Ruslan Salakhutdinov, and Christopher D Manning. 2018. Hotpotqa: A dataset for diverse, explainable multi-hop question answering. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 2369–2380.
- Will Yeadon and Douglas P Halliday. 2023. Exploring durham university physics exams with large language models. *arXiv preprint arXiv:2306.15609*.
- Tao Yu, Rui Zhang, Kai Yang, Michihiro Yasunaga, Dongxu Wang, Zifan Li, James Ma, Irene Li, Qingning Yao, Shanelle Roman, et al. 2018. Spider: A large-scale human-labeled dataset for complex and cross-domain semantic parsing and text-to-sql task. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 3911–3921.
- Weihao Yu, Zihang Jiang, Yanfei Dong, and Jiashi Feng. 2020. Reclor: A reading comprehension dataset requiring logical reasoning. *arXiv preprint arXiv:2002.04326*.
- Kexun Zhang, Danqing Wang, Jingtao Xia, William Yang Wang, and Lei Li. 2023a. Algo: Synthesizing algorithmic programs with generated oracle verifiers. *arXiv preprint arXiv:2305.14591*.
- Ziyin Zhang, Chaoyu Chen, Bingchang Liu, Cong Liao, Zi Gong, Hang Yu, Jianguo Li, and Rui Wang. 2023b. [A survey on language models for code](#).
- Qinkai Zheng, Xiao Xia, Xu Zou, Yuxiao Dong, Shan Wang, Yufei Xue, Zihan Wang, Lei Shen, Andi Wang, Yang Li, et al. 2023. Codegeex: A pre-trained model for code generation with multilingual evaluations on humaneval-x. *arXiv preprint arXiv:2303.17568*.
- Kun Zhou, Yutao Zhu, Zhipeng Chen, Wentong Chen, Wayne Xin Zhao, Xu Chen, Yankai Lin, Ji-Rong Wen, and Jiawei Han. 2023. Don’t make your llm an evaluation benchmark cheater. *arXiv preprint arXiv:2311.01964*.

## A Dataset Details

Tag	#Problems	Tag	#Problems
implementation	1746	greedy	1441
math	1382	brute force	825
constructive algorithms	783	dp	577
sortings	514	data structures	391
strings	381	binary search	342
number theory	309	graphs	263
dfs and similar	244	two pointers	197
combinatorics	179	bitmasks	154
geometry	142	trees	137
games	87	dsu	84
shortest paths	66	*special	58
probabilities	52	hashing	48
divide and conquer	35	flows	24
graph matchings	22	ternary search	22
matrices	22	expression parsing	19
string suffix structures	10	2-sat	7
chinese remainder theorem	5	schedules	4
meet-in-the-middle	4	fft	4

Table 5: Statistics of the types of problems in **D1**, **D2**, **D3**.

## B Experiment Details

Year	WA1	WA2	WA3	WA4+	RE	CE	TLE	MLE
2010	0.49	0.15	0.04	0.19	0.04	0.03	0.01	0.04
2011	0.53	0.13	0.10	0.06	0.07	0.06	0.04	0.01
2012	0.55	0.14	0.05	0.12	0.04	0.04	0.03	0.02
2013	0.63	0.13	0.06	0.07	0.03	0.03	0.04	0.01
2014	0.53	0.18	0.05	0.09	0.04	0.05	0.04	0.01
2015	0.62	0.10	0.02	0.14	0.03	0.02	0.04	0.02
2016	0.68	0.09	0.02	0.10	0.02	0.05	0.04	0.00
2017	0.66	0.13	0.05	0.08	0.03	0.03	0.03	0.02
2018	0.58	0.12	0.06	0.13	0.01	0.03	0.05	0.01
2019	0.67	0.11	0.04	0.06	0.02	0.05	0.02	0.02
2020	0.77	0.06	0.01	0.01	0.05	0.06	0.03	0.01
2021	0.82	0.07	0.02	0.01	0.02	0.04	0.02	0.00
2022	0.91	0.05	0.00	0.00	0.02	0.00	0.02	0.00
2023	0.90	0.03	0.00	0.00	0.03	0.01	0.02	0.01
Average	0.70	0.10	0.03	0.06	0.03	0.03	0.03	0.01

Table 6: Error category of GPT-4 from 2010 to 2023. The abbreviations stand for: WA1, WA2, WA3, and WA4+ (Wrong Answers on Test 1, 2, 3, and 4 or above), RE (Runtime Error), CE (Compilation Error), TLE (Time Limit Exceeded), and MLE (Memory Limit Exceeded).

## C Prompt Details

$\rho_1$	<p>You are given a problem, you need to write a C++ solution and explain the algorithm.</p> <p>{<i>promblem_name</i>}</p> <p>{<i>promblem_description</i>}</p> <p>Input specification: {<i>input_format</i>}</p> <p>Output specification: {<i>output_format</i>}</p> <p>Note: {<i>note</i>}</p> <p>Memory limit: {<i>memory_limit</i>}</p> <p>Time limit: {<i>time_limit</i>}</p> <p>Example:</p> <p>Input:</p> <p>{<i>input<sub>i</sub></i>}</p> <p>Output:</p> <p>{<i>output<sub>i</sub></i>}</p> <p>Please provide a C++ code in ```cpp\n...\n```</p>
$\rho_2$	<p>Read the problem, write a C++ solution and explain the algorithm. {<i>promblem_name</i>}:          {<i>promblem_description</i>} Input specification is {<i>input_format</i>}. Output specification is {<i>output_format</i>}. Note that {<i>note</i>}. Memory limit is {<i>memory_limit</i>}. Time limit is {<i>time_limit</i>}. Example <i>i</i> input is {<i>input<sub>i</sub></i>}. Example <i>i</i> output is {<i>output<sub>i</sub></i>}. Please provide a C++ code in ```cpp\n...\n```</p>
$\rho_3$	<p>Finish the solution of this programming problem.</p> <p>{<i>promblem_name</i>}</p> <p>{<i>promblem_description</i>}</p> <p>Input specification: {<i>input_format</i>}</p> <p>Output specification: {<i>output_format</i>}</p> <p>Note: {<i>note</i>}</p> <p>Memory limit: {<i>memory_limit</i>}</p> <p>Time limit: {<i>time_limit</i>}</p> <p>Example:</p> <p>Input:</p> <p>{<i>input<sub>i</sub></i>}</p> <p>Output:</p> <p>{<i>output<sub>i</sub></i>}</p> <p>C++ code solution:</p> <p>```cpp</p>
$\rho_{cot}$	<p>You are given an algorithm problem. First, provide a detailed explanation of the algorithm solution, including the logic behind it, the time and space complexity, and any important considerations or edge cases. Then, implement the solution in C++ code, ensuring it is clean, efficient, and well-commented.</p> <p>{<i>promblem_name</i>}</p> <p>{<i>promblem_description</i>}</p> <p>Input specification: {<i>input_format</i>}</p> <p>Output specification: {<i>output_format</i>}</p> <p>Note: {<i>note</i>}</p> <p>Memory limit: {<i>memory_limit</i>}</p> <p>Time limit: {<i>time_limit</i>}</p> <p>Example:</p> <p>Input:</p> <p>{<i>input<sub>i</sub></i>}</p> <p>Output:</p> <p>{<i>output<sub>i</sub></i>}</p> <p>Please provide a C++ code in ```cpp\n...\n```</p>
$\rho_{sip}$	<p>Please extract the essential components from this algorithm problem for a C++ solution, removing any superfluous narrative or context.</p> <p>{<i>promblem_name</i>}</p> <p>{<i>promblem_description</i>}</p> <p>Input specification: {<i>input_format</i>}</p> <p>Output specification: {<i>output_format</i>}</p> <p>Note: {<i>note</i>}</p>
$\rho_{sipgen}$	<p>You are given a problem, you need to write a C++ solution and explain the algorithm.</p> <p>{<i>promblem_simlified</i>}</p> <p>Memory limit: {<i>memory_limit</i>}</p> <p>Time limit: {<i>time_limit</i>}</p> <p>Example:</p> <p>Input:</p> <p>{<i>input<sub>i</sub></i>}</p> <p>Output:</p> <p>{<i>output<sub>i</sub></i>}</p> <p>Please provide a C++ code in ```cpp\n...\n```</p>

Table 7: Prompts used in this study.