

# linear\_regression\_interpretability

February 6, 2024

## 1 Interpretability of a linear regression model

Interpret means to “explain” or to present in understandable terms. The ability to express in understandable terms, what the model has learned and the reasons that affect their output.

Interpretability is about the extent to which a cause and effect can be observed within a system. Or to put it another way, it is the extent to which you are able to predict what is going to happen, given a change in input or algorithmic parameters. It’s being able to understand which inputs are the most predictive (i.e., impact the prediction/output the most), and anticipate how predictions will change with differing inputs.

- If a customer is rejected a loan, we can say why
- If an insurance provides a certain premium, we know the reasons.
- If we diagnose a patient with a certain disease, we can tell them why

### 1.1 Fit an interpretable linear regression model and make global and local interpretations

The idea is to fit an interpretable linear regression model, evaluate the model fit and the coefficients, and then interpret the predictions globally and locally.

In this example, we will use the shrinkage method and variable selection for linear regression models (called lasso regression) is another regularized version of linear regression.

An important peculiarity of lasso regression is that it tends to suppress the weights of less important features (i.e., set them to zero). Roughly speaking, lasso regression automatically performs feature selection and outputs a sparse model (i.e., with few non-zero feature weights).

The workflow is the following:

- Make some data engineer to prepare the data
- Exploratory data analysis and identify multi-collinearity
- Fit a linear model with the highest performance and least number of features
- Evaluate the model fit
- Evaluate the coefficients (global interpretation)
- Evaluate a few observations individually (local interpretation)

```
[1]: # imports
import warnings
warnings.filterwarnings("ignore")
```

```

import math
import numpy as np
import pandas as pd
pd.set_option('display.max_columns', 10000)
# pd.set_option('display.max_rows', 10000)
import matplotlib.pyplot as plt
import seaborn as sns
import sweetviz as sv
from itertools import product
from scipy import stats
import statsmodels.api as sm

import scipy.stats as ss
from scipy.stats.contingency import association
from statsmodels.stats.outliers_influence import variance_inflation_factor

from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.pipeline import Pipeline, make_pipeline
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
from sklearn.linear_model import Lasso
from sklearn.model_selection import cross_validate
from sklearn.base import BaseEstimator, TransformerMixin
from sklearn.model_selection import GridSearchCV, train_test_split

```

```

[2]: # load dataset
train_set = pd.read_csv('datasets/house_price_train.csv')
train_set.head()

```

```

[2]:
   Id  MSSubClass  MSZoning  LotFrontage  LotArea  Street  Alley  LotShape  \
0    1           60        RL           65.0    8450   Pave   NaN      Reg
1    2           20        RL           80.0    9600   Pave   NaN      Reg
2    3           60        RL           68.0   11250   Pave   NaN      IR1
3    4           70        RL           60.0    9550   Pave   NaN      IR1
4    5           60        RL           84.0   14260   Pave   NaN      IR1

   LandContour  Utilities  LotConfig  LandSlope  Neighborhood  Condition1  \
0          Lvl1    AllPub    Inside      Gtl      CollgCr      Norm
1          Lvl1    AllPub      FR2      Gtl      Veenker      Feedr
2          Lvl1    AllPub    Inside      Gtl      CollgCr      Norm
3          Lvl1    AllPub    Corner      Gtl      Crawfor      Norm
4          Lvl1    AllPub      FR2      Gtl      NoRidge      Norm

   Condition2  BldgType  HouseStyle  OverallQual  OverallCond  YearBuilt  \
0          Norm     1Fam     2Story           7            5        2003
1          Norm     1Fam     1Story           6            8        1976
2          Norm     1Fam     2Story           7            5        2001

```

3	Norm	1Fam	2Story	7	5	1915
4	Norm	1Fam	2Story	8	5	2000

	YearRemodAdd	RoofStyle	RoofMatl	Exterior1st	Exterior2nd	MasVnrType	\
0	2003	Gable	CompShg	VinylSd	VinylSd	BrkFace	
1	1976	Gable	CompShg	MetalSd	MetalSd	NaN	
2	2002	Gable	CompShg	VinylSd	VinylSd	BrkFace	
3	1970	Gable	CompShg	Wd Sdng	Wd Shng	NaN	
4	2000	Gable	CompShg	VinylSd	VinylSd	BrkFace	

	MasVnrArea	ExterQual	ExterCond	Foundation	BsmtQual	BsmtCond	BsmtExposure	\
0	196.0	Gd	TA	PConc	Gd	TA	No	
1	0.0	TA	TA	CBlock	Gd	TA	Gd	
2	162.0	Gd	TA	PConc	Gd	TA	Mn	
3	0.0	TA	TA	BrkTil	TA	Gd	No	
4	350.0	Gd	TA	PConc	Gd	TA	Av	

	BsmtFinType1	BsmtFinSF1	BsmtFinType2	BsmtFinSF2	BsmtUnfSF	TotalBsmtSF	\
0	GLQ	706	Unf	0	150	856	
1	ALQ	978	Unf	0	284	1262	
2	GLQ	486	Unf	0	434	920	
3	ALQ	216	Unf	0	540	756	
4	GLQ	655	Unf	0	490	1145	

	Heating	HeatingQC	CentralAir	Electrical	1stFlrSF	2ndFlrSF	LowQualFinSF	\
0	GasA	Ex	Y	SBrkr	856	854	0	
1	GasA	Ex	Y	SBrkr	1262	0	0	
2	GasA	Ex	Y	SBrkr	920	866	0	
3	GasA	Gd	Y	SBrkr	961	756	0	
4	GasA	Ex	Y	SBrkr	1145	1053	0	

	GrLivArea	BsmtFullBath	BsmtHalfBath	FullBath	HalfBath	BedroomAbvGr	\
0	1710	1	0	2	1	3	
1	1262	0	1	2	0	3	
2	1786	1	0	2	1	3	
3	1717	1	0	1	0	3	
4	2198	1	0	2	1	4	

	KitchenAbvGr	KitchenQual	TotRmsAbvGrd	Functional	Fireplaces	FireplaceQu	\
0	1	Gd	8	Typ	0	NaN	
1	1	TA	6	Typ	1	TA	
2	1	Gd	6	Typ	1	TA	
3	1	Gd	7	Typ	1	Gd	
4	1	Gd	9	Typ	1	TA	

	GarageType	GarageYrBlt	GarageFinish	GarageCars	GarageArea	GarageQual	\
0	Attchd	2003.0	RFn	2	548	TA	

1	Attchd	1976.0	RFn	2	460	TA
2	Attchd	2001.0	RFn	2	608	TA
3	Detchd	1998.0	Unf	3	642	TA
4	Attchd	2000.0	RFn	3	836	TA

	GarageCond	PavedDrive	WoodDeckSF	OpenPorchSF	EnclosedPorch	3SsnPorch	\
0	TA	Y	0	61	0	0	
1	TA	Y	298	0	0	0	
2	TA	Y	0	42	0	0	
3	TA	Y	0	35	272	0	
4	TA	Y	192	84	0	0	

	ScreenPorch	PoolArea	PoolQC	Fence	MiscFeature	MiscVal	MoSold	YrSold	\
0	0	0	NaN	NaN	NaN	0	2	2008	
1	0	0	NaN	NaN	NaN	0	5	2007	
2	0	0	NaN	NaN	NaN	0	9	2008	
3	0	0	NaN	NaN	NaN	0	2	2006	
4	0	0	NaN	NaN	NaN	0	12	2008	

	SaleType	SaleCondition	SalePrice
0	WD	Normal	208500
1	WD	Normal	181500
2	WD	Normal	223500
3	WD	Abnorml	140000
4	WD	Normal	250000

## 1.2 Exploratory data analysis

This topic we gonna work in an exploration to see what we should do with this data to be able to go to the next steps, readers can skip this step if you are interested just in the model interpretation step.

[Link to the dataset](#)

Obs: we don't gonna make an extensive and deep exploratory, because the goal of this notebook is to show how to interpret the models, but in a real project, you should go deeper in the exploration and extract many information as possible.

### 1.2.1 Univariate data analysis

For this step, we gonna use a very nice tool to make the things faster that is [sweetviz](#) tool. If you don't know the tool, have a look in the documentation!

```
[3]: # general informations about the dataset
train_set.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 81 columns):
```

#	Column	Non-Null Count	Dtype
0	Id	1460 non-null	int64
1	MSSubClass	1460 non-null	int64
2	MSZoning	1460 non-null	object
3	LotFrontage	1201 non-null	float64
4	LotArea	1460 non-null	int64
5	Street	1460 non-null	object
6	Alley	91 non-null	object
7	LotShape	1460 non-null	object
8	LandContour	1460 non-null	object
9	Utilities	1460 non-null	object
10	LotConfig	1460 non-null	object
11	LandSlope	1460 non-null	object
12	Neighborhood	1460 non-null	object
13	Condition1	1460 non-null	object
14	Condition2	1460 non-null	object
15	BldgType	1460 non-null	object
16	HouseStyle	1460 non-null	object
17	OverallQual	1460 non-null	int64
18	OverallCond	1460 non-null	int64
19	YearBuilt	1460 non-null	int64
20	YearRemodAdd	1460 non-null	int64
21	RoofStyle	1460 non-null	object
22	RoofMatl	1460 non-null	object
23	Exterior1st	1460 non-null	object
24	Exterior2nd	1460 non-null	object
25	MasVnrType	588 non-null	object
26	MasVnrArea	1452 non-null	float64
27	ExterQual	1460 non-null	object
28	ExterCond	1460 non-null	object
29	Foundation	1460 non-null	object
30	BsmtQual	1423 non-null	object
31	BsmtCond	1423 non-null	object
32	BsmtExposure	1422 non-null	object
33	BsmtFinType1	1423 non-null	object
34	BsmtFinSF1	1460 non-null	int64
35	BsmtFinType2	1422 non-null	object
36	BsmtFinSF2	1460 non-null	int64
37	BsmtUnfSF	1460 non-null	int64
38	TotalBsmtSF	1460 non-null	int64
39	Heating	1460 non-null	object
40	HeatingQC	1460 non-null	object
41	CentralAir	1460 non-null	object
42	Electrical	1459 non-null	object
43	1stFlrSF	1460 non-null	int64
44	2ndFlrSF	1460 non-null	int64
45	LowQualFinSF	1460 non-null	int64

```

46  GrLivArea      1460 non-null  int64
47  BsmtFullBath   1460 non-null  int64
48  BsmtHalfBath   1460 non-null  int64
49  FullBath        1460 non-null  int64
50  HalfBath        1460 non-null  int64
51  BedroomAbvGr   1460 non-null  int64
52  KitchenAbvGr   1460 non-null  int64
53  KitchenQual     1460 non-null  object
54  TotRmsAbvGrd   1460 non-null  int64
55  Functional      1460 non-null  object
56  Fireplaces      1460 non-null  int64
57  FireplaceQu     770 non-null   object
58  GarageType      1379 non-null  object
59  GarageYrBlt     1379 non-null  float64
60  GarageFinish    1379 non-null  object
61  GarageCars      1460 non-null  int64
62  GarageArea      1460 non-null  int64
63  GarageQual      1379 non-null  object
64  GarageCond      1379 non-null  object
65  PavedDrive      1460 non-null  object
66  WoodDeckSF      1460 non-null  int64
67  OpenPorchSF     1460 non-null  int64
68  EnclosedPorch   1460 non-null  int64
69  3SsnPorch       1460 non-null  int64
70  ScreenPorch     1460 non-null  int64
71  PoolArea        1460 non-null  int64
72  PoolQC          7 non-null     object
73  Fence           281 non-null   object
74  MiscFeature     54 non-null     object
75  MiscVal         1460 non-null  int64
76  MoSold          1460 non-null  int64
77  YrSold          1460 non-null  int64
78  SaleType        1460 non-null  object
79  SaleCondition   1460 non-null  object
80  SalePrice       1460 non-null  int64
dtypes: float64(3), int64(35), object(43)
memory usage: 924.0+ KB

```

```

[4]: my_report = sv.analyze(train_set, 'SalePrice')
      my_report.show_html() # Default arguments will generate to "SWEETVIZ_REPORT.
                             ↪html"

```

```

Feature: SalePrice (TARGET) | | [ 1%] 00:00 ->
(00:02 left)

Feature: MSZoning | | [ 5%] 00:02 ->
(00:45 left)

```

### 1.2.2 Bivariate Data Analysis

Here we will deal with the descriptive analysis of the **association** between two variables. **In general, we say that there is an association between two variables if knowledge of the value of one of them gives us some information about some characteristic of the distribution (of frequencies) of the other.**

*We can highlight three cases:*

1. both variables are qualitative.
2. both variables are quantitative.
3. one variable is qualitative and the other is quantitative.

**Two qualitative variables and evaluate multicollinearity** Here the idea is to check the correlation between two qualitative variables.

What are the consequences of multicollinearity?

If two variables are perfectly collinear, in other words, if they have correlation coefficient equal to 1, then what happens is that there is an infinite combination of coefficients (betas) that would work equally well. So basically we have an infinite number of linear regression models that will predict equally well the target from these two perfectly collinear variables. Which means that we are not able to understand what is the real relationship between those variables and the target.

- Perfect collinearity is rare
- Partial collinearity is unavoidable

So what happens is that when we have correlated variables, one of the terms (feature x coefficient) will account for a degree of the variability, and then the other term basically accounts for the remaining variability that is not explained, but in both cases the coefficient doesn't really represent the real if you want association between the variable and the target.

```
[3]: # filter the qualitative variables
qualitative_features_encoding = [
    'MSZoning', 'Street', 'LotShape', 'LandContour',
    'Utilities', 'LotConfig', 'LandSlope', 'Neighborhood',
    'Condition1', 'Condition2', 'BldgType', 'HouseStyle',
    'RoofStyle', 'RoofMat1', 'Exterior1st', 'Exterior2nd',
    'Foundation', 'Heating', 'CentralAir', 'Functional',
    'PavedDrive', 'SaleType', 'SaleCondition']

qualitative_features_missing_mode = [
    'Electrical', 'GarageType']

ordinal_features_quality = [
    'ExterQual', 'ExterCond',
    'BsmtQual', 'BsmtCond',
    'HeatingQC', 'KitchenQual',
    'GarageQual', 'GarageCond']

ordinal_features_exposure = ['BsmtExposure']
```

```

ordinal_features_finish = ['BsmtFinType1', 'BsmtFinType2']

ordinal_features_garage = ['GarageFinish']

qualitative_vars = qualitative_features_encoding +
    ↳ qualitative_features_missing_mode + ordinal_features_quality +
    ↳ ordinal_features_exposure + ordinal_features_finish + ordinal_features_garage

```

```

[4]: # create a dataframe with only categorical variables
categorical_df = train_set[qualitative_vars]

# removing records with at least one null value in a row
df_cat_v1 = categorical_df.dropna()

## let us split this list into two parts
cat_var1 = qualitative_vars
cat_var2 = qualitative_vars

# let us jump to Chi-Square test
# creating all possible combinations between the above two variables list
cat_var_prod = list(product(cat_var1, cat_var2, repeat = 1))

# creating an empty variable and picking only the p value from the output of
↳ Chi-Square test
result = []
for i in cat_var_prod:
    if i[0] != i[1]:
        contingency_table = pd.crosstab(df_cat_v1[i[0]], df_cat_v1[i[1]])
        chi2_pval = ss.chi2_contingency(contingency_table)[1]
        tschuprow_pval = association(contingency_table, method='tschuprow')
        result.append((i[0], i[1], chi2_pval, tschuprow_pval))

# Creating dataframe
result_df = pd.DataFrame(result, columns=['Variable_1', 'Variable_2',
    ↳ 'Chi2_P_Value', 'Tschuprow'])

# let's filter the values with tschuprow coefficient higher than 0.7 to catch
↳ multicollinearity
result_df.loc[(result_df['Tschuprow'] >= 0.7)]

```

```

[4]:
Variable_1  Variable_2  Chi2_P_Value  Tschuprow
518  Exterior1st  Exterior2nd         0.0    0.744531
554  Exterior2nd  Exterior1st         0.0    0.744531

```

Maybe we can remove one of these two features, the one that has less correlation with the target is a good choice.



**Two quantitative variables and evaluate multicollinearity** Here the idea is to check the correlation between two quantitative variables and check multicollinearity.


```
[5]: # select features according to their types and missing values
quantitative_features = ['MSSubClass', 'LotArea', 'OverallQual', 'OverallCond',
                          'YearBuilt', 'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF',
                          'TotalBsmtSF', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF',
                          'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath',
                          'FullBath',
                          'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr',
                          'TotRmsAbvGrd',
                          'Fireplaces', 'GarageCars', 'GarageArea', 'WoodDeckSF',
                          'OpenPorchSF', 'EnclosedPorch', '3SsnPorch',
                          'ScreenPorch',
                          'PoolArea', 'MiscVal', 'MoSold', 'YrSold']

quantitative_features_missing_median = ['LotFrontage', 'MasVnrArea']

quantitative_features_missing_mode = ['LotFrontage', 'YearRemodAdd',
                                       'GarageYrBlt']

quantitative_vars = quantitative_features +
                    quantitative_features_missing_median + quantitative_features_missing_mode

[6]: # quantitative associations
quantitative_df = train_set[quantitative_vars]
corr = quantitative_df.corr()
plt.figure(figsize=(15,15))
sns.heatmap(corr, fmt='.0%', annot=True, square=True)
plt.tight_layout()
```



From the Pearson correlation, few variables appeared to be highly correlated. Let's check this collinearity with the VIF metric and that could possibly be an indication for feature selection.

The second metric for gauging multicollinearity is the variance inflation factor (VIF). The VIF directly measures the ratio of the variance of the entire model to the variance of a model with only the feature in question.

In layman's terms, it gauges how much a feature's inclusion contributes to the overall variance of the coefficients of the features

values is likely to be contributing to multicollinearity.

```
[7]: # Compute VIF data for each independent variable
clean_quantitative_df = quantitative_df.dropna()
vif = pd.DataFrame()
vif["features"] = clean_quantitative_df.columns
vif["vif_Factor"] = [variance_inflation_factor(clean_quantitative_df.values, i)
    ↪for i in range(clean_quantitative_df.shape[1])]
vif.sort_values(by=['vif_Factor'], ascending=False, inplace=True)
vif
```

```
[7]:
```

	features	vif_Factor
8	TotalBsmtSF	inf
10	2ndFlrSF	inf
32	LotFrontage	inf
34	LotFrontage	inf
12	GrLivArea	inf
5	BsmtFinSF1	inf
6	BsmtFinSF2	inf
7	BsmtUnfSF	inf
11	LowQualFinSF	inf
9	1stFlrSF	inf
36	GarageYrBlt	2.680925e+04
31	YrSold	2.528962e+04
4	YearBuilt	2.466581e+04
35	YearRemodAdd	2.450464e+04
19	TotRmsAbvGrd	8.364791e+01
2	OverallQual	7.308071e+01
3	OverallCond	4.964269e+01
18	KitchenAbvGr	4.203815e+01
21	GarageCars	3.980117e+01
22	GarageArea	3.517381e+01
17	BedroomAbvGr	3.381041e+01
15	FullBath	2.874419e+01
30	MoSold	6.840270e+00
0	MSSubClass	4.821513e+00
13	BsmtFullBath	3.650975e+00
16	HalfBath	3.594106e+00
1	LotArea	3.451820e+00
20	Fireplaces	3.028183e+00
24	OpenPorchSF	1.965871e+00
23	WoodDeckSF	1.944755e+00
33	MasVnrArea	1.937282e+00
25	EnclosedPorch	1.486659e+00
27	ScreenPorch	1.239925e+00
14	BsmtHalfBath	1.214943e+00
28	PoolArea	1.192611e+00

```
29         MiscVal  1.119622e+00
26         3SsnPorch 1.043918e+00
```

We see that a lot of variables have a high VIF and therefore, it may be variables to eliminate in the modeling that we will do later.

**One qualitative and one quantitative variable** Here the idea is to check the correlation between the qualitative and quantitative variable (the target) to see if the qualitative variables have a high influence in the target, but for this project we gonna skip this part, but in your project you should go deeper.

### 1.3 Preprocessing

Based on what we have seen in exploratory data analysis, we gonna make some transformations in the data to be able to fit them in the linear regression model.

```
[8]: # select only the features that we are going to use
X = train_set.drop(['SalePrice'], axis=1)
y = train_set['SalePrice']
```

```
[9]: class QualMapper(BaseEstimator, TransformerMixin):
    def __init__(self, qual_vars):
        self.qual_vars = qual_vars

    def fit(self, X, y=None):
        return self # no need to do anything here

    def transform(self, X):
        def map_quality(entry):
            if entry == 'Po':
                return 1
            elif entry == 'Fa':
                return 2
            elif entry == 'TA':
                return 3
            elif entry == 'Gd':
                return 4
            elif entry == 'Ex':
                return 5
            else:
                return 0 # or 'Missing' if you prefer to keep it as a string

        for var in self.qual_vars:
            X[var] = X[var].fillna('Missing')
            X[var] = X[var].apply(map_quality)

        return X
```

```

def get_feature_names_out(self, input_features=None):
    return input_features

class ExposureMapper(BaseEstimator, TransformerMixin):
    def __init__(self, expo_vars):
        self.expo_vars = expo_vars

    def fit(self, X, y=None):
        return self # no need to do anything here

    def transform(self, X):
        def map_expo(entry):
            if entry == 'No':
                return 1
            elif entry == 'Mn':
                return 2
            elif entry == 'Av':
                return 3
            elif entry == 'Gd':
                return 4
            else:
                return 0 # or 'Missing' if you prefer to keep it as a string

        for var in self.expo_vars:
            X[var] = X[var].fillna('Missing')
            X[var] = X[var].apply(map_expo)

        return X

    def get_feature_names_out(self, input_features=None):
        return input_features

class FinishMapper(BaseEstimator, TransformerMixin):
    def __init__(self, finish_vars):
        self.finish_vars = finish_vars

    def fit(self, X, y=None):
        return self # no need to do anything here

    def transform(self, X):
        def map_finish(entry):
            if entry == 'Unf':
                return 1
            elif entry == 'LwQ':
                return 2

```

```

        elif entry == 'Rec':
            return 3
        elif entry == 'BLQ':
            return 4
        elif entry == 'ALQ':
            return 5
        elif entry == 'GLQ':
            return 6
        else:
            return 0 # or 'Missing' if you prefer to keep it as a string

    for var in self.finish_vars:
        X[var] = X[var].fillna('Missing')
        X[var] = X[var].apply(map_finish)

    return X

def get_feature_names_out(self, input_features=None):
    return input_features

class GarageMapper(BaseEstimator, TransformerMixin):
    def __init__(self, garage_vars):
        self.garage_vars = garage_vars

    def fit(self, X, y=None):
        return self # no need to do anything here

    def transform(self, X):
        def map_garage(entry):
            if entry == 'Unf':
                return 1
            elif entry == 'RFn':
                return 2
            elif entry == 'Fin':
                return 3
            else:
                return 0 # or 'Missing' if you prefer to keep it as a string

        for var in self.garage_vars:
            X[var] = X[var].fillna('Missing')
            X[var] = X[var].apply(map_garage)

        return X

    def get_feature_names_out(self, input_features=None):
        return input_features

```

```

# class RemoveRareCategories(BaseEstimator, TransformerMixin):
#     def __init__(self, quali_vars):
#         self.quali_vars = quali_vars

#     def fit(self, X, y=None):
#         return self # no need to do anything here

#     def transform(self, X):
#         def map_rare_entries(series):
#             frequency_entries = series.groupby(series).transform('count') /
#             len(series)
#             rare_entries = frequency_entries < 0.01
#             print(rare_entries)
#             return np.where(rare_entries, 'Rare', series)

#         for var in self.quali_vars:
#             X[var] = map_rare_entries(X[var])

#         return X

```

```
[10]: processed_features = quantitative_vars + qualitative_vars
```

```

quantitative_preproc = make_pipeline(
    StandardScaler())

quantitative_median_preproc = make_pipeline(
    SimpleImputer(strategy='median'),
    StandardScaler())

quantitative_mode_preproc = make_pipeline(
    SimpleImputer(strategy='most_frequent'),
    StandardScaler())

qualitative_preproc = make_pipeline(
    OneHotEncoder(handle_unknown='ignore'))

qualitative_mode_preproc = make_pipeline(
    OneHotEncoder(handle_unknown='ignore'))

ordinal_quality_preproc = make_pipeline(
    QualMapper(ordinal_features_quality),
    StandardScaler())

ordinal_exposure_preproc = make_pipeline(
    ExposureMapper(ordinal_features_exposure),

```

```

StandardScaler())

ordinal_finish_preproc = make_pipeline(
    FinishMapper(ordinal_features_finish),
    StandardScaler())

ordinal_garage_preproc = make_pipeline(
    GarageMapper(ordinal_features_garage),
    StandardScaler())

# apply the respective transformations with columntransformer method
preprocessor = ColumnTransformer([
    ('quantitative_preproc', quantitative_preproc, quantitative_features),
    ('quantitative_median_preproc', quantitative_median_preproc,
↳quantitative_features_missing_median),
    ('quantitative_mode_preproc', quantitative_mode_preproc,
↳quantitative_features_missing_mode),
    ('qualitative_preproc', qualitative_preproc,
↳qualitative_features_encoding),
    ('qualitative_mode_preproc', qualitative_mode_preproc,
↳qualitative_features_missing_mode),
    ('ordinal_quality_preproc', ordinal_quality_preproc,
↳ordinal_features_quality),
    ('ordinal_exposure_preproc', ordinal_exposure_preproc,
↳ordinal_features_exposure),
    ('ordinal_finish_preproc', ordinal_finish_preproc, ordinal_features_finish),
    ('ordinal_garage_preproc', ordinal_garage_preproc,
↳ordinal_features_garage)],
    remainder='drop')

```

## 1.4 Fit an interpretable linear model

We'll start by understanding how linear regression model works, then we'll move on to learning statistical tests that we can use to evaluate the performance of these models, and this is important because if the model doesn't fit the data well then we cannot extract meaningful interpretations, and then we'll move on to interpret the output of the linear regression model (Lasso) at a global and local level.

Using Lasso, train the model that performs the best and has the least number of features.

If you get errors, reduce the penalization.

```

[11]: def run_regressor_models(X, y, cv, scoring):
    '''Function that trains the following machine learning models:
    DecisionTreeRegressor, RandomForestRegressor, svm, SGDRegressor,
    GradientBoostingRegressor, ElasticNet, MLPRegressor.
    The function applies cross-validation on the dataset and returns the mean

```



and standard deviation of the selected metric on the training and validation set.

The only active metrics are RMSE and R2.

:param X: (dataframe or numpy array)  
DataFrame or array with the set of independent variables.

:param y: (series or numpy array)  
Column or array with the dependent variable.

:param cv: (int)  
Determines the cross-validation splitting strategy.

:param scoring: (str)  
Strategy to evaluate the performance of the cross-validated model on the validation set.  
Should be passed within quotes when calling the function.

:return: (dataframe)  
DataFrame with models, mean, and standard deviation on training and validation set.

```
'''
# Instantiate the models
reg = Pipeline(
    steps=[('preprocessor', preprocessor),
           ('regressor', Lasso(random_state=42))]
)
scores = cross_validate(reg, X, y, return_train_score=True,
                        scoring=scoring, cv=cv, return_estimator=True)

# Train and test RMSE
if scoring == 'neg_mean_squared_error':
    train_rmse_scores = np.sqrt(-scores['train_score'])
    test_rmse_scores = np.sqrt(-scores['test_score'])
    mean_train = train_rmse_scores.mean()
    mean_test = test_rmse_scores.mean()
    std_train = train_rmse_scores.std()
    std_test = test_rmse_scores.std()

# Train and test R2
if scoring == 'r2':
    train_r2_scores = scores['train_score']
    test_r2_scores = scores['test_score']
    mean_train = train_r2_scores.mean()
    mean_test = test_r2_scores.mean()
    std_train = train_r2_scores.std()
    std_test = test_r2_scores.std()
```

```

# Create final dataset
df_result = pd.DataFrame(
    {'MODEL': reg[1], 'MEAN_TRAIN_SCORES': mean_train,
     'MEAN_TEST_SCORES': mean_test, 'STD_TRAIN_SCORES': std_train,
     'STD_TEST_SCORES': std_test}, index=[0])

return df_result

```

```

[12]: df_result = run_regressor_models(X, y, 5, 'r2')
df_result

```

```

[12]:
          MODEL  MEAN_TRAIN_SCORES  MEAN_TEST_SCORES  \
0  Lasso(random_state=42)          0.923311          0.812339

          STD_TRAIN_SCORES  STD_TEST_SCORES
0              0.003282          0.07568

```

As we can see, we have some overfit, probably because we have a lot of features and a small number of instances. Let's try to decrease this overfitting increasing the regularization parameter.

In this cases when the number of features are big and we don't have a big number of rows, it's better to use the regularized models like Lasso. The term regularization refers to a set of techniques used to specify models that fit a set of data while avoiding overfitting. In general terms, these techniques serve to fit regression models based on a cost function that contains a penalty term. This term is intended to reduce the influence of coefficients responsible for excessive fluctuations.

When there are predictor variables that are not associated with the response variable, the regression models adjusted by least squares (like LinearRegression from sklearn) may be more complex than desired, as the coefficients associated with these variables will not be canceled.

## 1.5 Fit, evaluate and tune model in test set

To evaluate the model, we often use a metric that is called  $R^2$ . To summary, the  $R^2$  is the fraction of variability explained by the model. Independent of the result of  $R^2$ , how do we know that this value is statistically significant?

So, to evaluate the  $R^2$  and that fit with confidence, we need statistical tests. The statistical tests that we gonna use are f-statistic that basically capture the relationship between the variance that is explained by the model versus the variance that is not explained by the model. In this case in we want a big f-statistic (the model explain more variability) and a small p-value. If this is the case, then the model offers a good fit of the data and the interpretations that we derive from it are meaningful.

F-statistic follows a known probability distribution for situations where a model is not a good fit and we gonna make a hypothesis test.

- Null hypothesis -> the model is not a good fit.
- Alternative hypothesis -> the model is a good fit.

```
[13]: # Hyperparameter tuning
# 1. Instantiate the pipeline
final_model = Pipeline(
    steps=[
        ('preprocessor', preprocessor),
        ('regressor', Lasso(random_state=42))
    ]
)

# 2. Hyperparameter interval to be tested
param_grid = {
    'regressor__alpha': [
#         0.01, 0.04, 0.1, 0.5, 1.0, 10.0,
#         100.0, 200.0, 300.0, 500.0, 1000.0,
        2500.0],
} # you should try as many values as possible, but to illustrate we gonna put a
    ↪heavy weight

# 3. Training and apply grid search with cross validation
grid_search = GridSearchCV(final_model, param_grid, cv = 5, scoring = 'r2',
                           return_train_score = True)
grid_search.fit(X, y)

# Seeing the best hyperparameters for the model
print('The best hyperparameters were:', grid_search.best_params_)
```

The best hyperparameters were: {'regressor\_\_alpha': 2500.0}

```
[14]: cvres = grid_search.cv_results_
cvres = [(mean_test_score,
         mean_train_score) for mean_test_score,
         mean_train_score in sorted(zip(cvres['mean_test_score'],
         cvres['mean_train_score']),
         reverse=True) if (math.
    ↪isnan(mean_test_score) != True)]
print(
    'The mean test score and mean train score is, respectively:',
    cvres[0])
```

The mean test score and mean train score is, respectively: (0.7999223632635879, 0.8202277519210949)

Overfitting continues, but we have a good model with  $R^2$  of 0.8 in test set. In other words, the model explains approximately 80% of the relationship between the dependent and independent variables.

The coefficient of determination must be accompanied by other tools for assessing fit, as it is not aimed at identifying whether all model assumptions are compatible with the data under investigation. Some of these tools are: residual graphs, cook graphs and local influence graphs. Here, we

will only talk about the residual graph.

## 1.6 Assumptions that must be met in a linear regression model

In summary, there are 5 basic assumptions of the regression algorithm that would be interesting to meet:

- Linear relationship between the independent variables and the target (in the case of linear regression)
- Little or no multicollinearity between variables
- Homoscedasticity assumption
- Normal distribution of error terms (if you want to test hypotheses about the model coefficients or construct confidence intervals for them)
- Little or no autocorrelation in residuals


When these assumptions are not met, regression analysis results can be misleading and the model may not perform well.

Here, we will calculate Homoscedasticity and the normal distribution of errors, since the other items have already been checked.

```
[15]: y_pred = grid_search.predict(X)
      residuals = y - y_pred

      # Standardize residuals
      residuals_standardized = residuals / residuals.std()

      # Residuals vs. Predicted Values Plot
      plt.scatter(y_pred, residuals_standardized)
      plt.xlabel("Predicted Values")
      plt.ylabel("Standardized Residuals")
      plt.title("Residuals vs. Predicted Values Plot (Standardized)")
      plt.axhline(y=0, color='r', linestyle='-') # Adding a line at y=0
      plt.show()
```



From the graph it appears that the residues are following some systematic pattern, which is not good. This suggests the presence of heteroscedasticity, that is, variances that are not constant over time. Outliers can also contribute to heteroscedasticity.


Now, let's check to the normality of errors.


```
[16]: # Checking for Normality of Errors
# Example using the Shapiro-Wilk test
_, p_value_sw = stats.shapiro(residuals_standardized)
if p_value_sw > 0.05:
    print("The residuals follow a normal distribution.")
else:
    print("The residuals do not follow a normal distribution.")

# Histogram of Residuals
plt.hist(residuals_standardized, bins=20)
plt.xlabel("Residuals")
plt.ylabel("Frequency")
plt.title("Histogram of Residuals")
plt.show()
```

```
# Q-Q Plot of Residuals  
stats.probplot(residuals_standardized, dist="norm", plot=plt)  
plt.title("Q-Q Plot of Residuals")  
plt.show()
```

The residuals do not follow a normal distribution.





Whenever we run a model or perform data analysis, it is common to check the distribution of the variables in question. If some are skewed and not normally distributed, we tend to worry. The truth is that the need to assume normality for independent and dependent variables is not always true.

**The variable that should be normally distributed is just the prediction error!**

The prediction error must follow a normal distribution with mean 0. The calculation of the confidence interval and the significance of the variable is based on this assumption.

You can also use the Shapiro-Wilk test to test the normality assumption. Basically, if the p-value is greater than the threshold (usually 0.05) then we accept the distribution as normal.

How do we fix the normality problem? Typically, there are 2 reasons why this problem occurs:

- Dependent or independent variables are very abnormal (you can see by the asymmetry and/or kurtosis of the variable)
- Existence of some outliers that hinder the model's prediction

We must then check for outliers in the variables and if this does not solve the problem, we must transform some variables that are not normally distributed so that they are normally distributed.

In conclusion, if you try to find a significant predictive factor or define the confidence interval, remember to check the distribution of the error term after building the model. If the dependent

variables or independent variables are very non-normal, one can use the box-cox transformation (for example) to transform it and make the error term more normally distributed.

In our case, neither homoscedasticity or the normal distribution of errors were met, which compromises our analysis. In your projects you must fix this, as this work is only at a didactic level, we will not delve into this issue further. One idea to fix this would perhaps be to normalize the target variable.

## 1.7 Evaluate the model globally

Linear regression models are intrinsically explainable, which means that if we understand how the model works, then we can make sense of their predictions.

Let's now try to interpret the model globally. For this we need to determine:

- Coefficient magnitude and sign
- Coefficient significance (t statistic and p-value)
- Effects plot

Determine the coefficient's error using cross-validation.

```
[17]: # Function to fit the model with GridSearchCV and return the coefficients
def fit_model_and_get_coeffs(X_train, y_train, param_grid):
    grid_search_globally = GridSearchCV(final_model, param_grid, cv=5,
    ↪scoring='r2', return_train_score=True)
    grid_search_globally.fit(X_train, y_train)
    return grid_search_globally.best_estimator_.named_steps.regressor.coef_,
    ↪grid_search_globally.best_estimator_.named_steps.preprocessor.
    ↪get_feature_names_out()

# Number of Bootstrap iterations
num_bootstrap_iterations = 100

# List to store coefficients for each Bootstrap iteration
coeffs_bootstrap = []

# Initialize list of coefficient names
coeffs_names = []

# Loop over each Bootstrap iteration
for i in range(num_bootstrap_iterations):
    # Random sampling with replacement of indices
    indices = np.random.choice(range(len(X)), size=len(X), replace=True)
    X_bootstrap, y_bootstrap = X.iloc[indices], y.iloc[indices]

    # Split the sampled dataset into train and test
    X_train, X_test, y_train, y_test = train_test_split(X_bootstrap,
    ↪y_bootstrap, test_size=0.2, random_state=42)

    # Fit the model and obtain coefficients for this sample
```



```

    coeffs_boot, new_coeffs_names = fit_model_and_get_coeffs(X_train, y_train,
↳param_grid)

    # Add unique coefficient names
    coeffs_names.extend([name for name in new_coeffs_names if name not in
↳coeffs_names])

    # Add coefficients of this Bootstrap iteration to the list
    coeffs_bootstrap.append(coeffs_boot)

# Determine the maximum length of coefficients array
max_coeffs_length = max(len(coeffs) for coeffs in coeffs_bootstrap)

# Fill coefficients arrays with zeros so that they all have the same length
coeffs_bootstrap_padded = [np.pad(coeffs, (0, max_coeffs_length - len(coeffs)),
↳mode='constant') for coeffs in coeffs_bootstrap]

# Convert the list of coefficients arrays into a numpy matrix
coeffs_bootstrap = np.vstack(coeffs_bootstrap_padded)

# Ensure coeffs_names has the same length as the number of columns in
↳coeffs_mean
coeffs_names = coeffs_names[:coeffs_bootstrap.shape[1]]

# Calculate the mean of coefficients
coeffs_mean = np.mean(coeffs_bootstrap, axis=0)

# Calculate the standard deviation of coefficients
coeffs_std = np.std(coeffs_bootstrap, axis=0)

# Create a DataFrame to store the mean and standard deviation of coefficients
coeffs_df = pd.DataFrame({'mean_coeffs_sign': coeffs_mean, 'coeffs_std':
↳coeffs_std}, index=coeffs_names).reset_index()
coeffs_df.sort_values(by=['mean_coeffs_sign'], ascending=False, inplace=True)
coeffs_df.rename(columns={'index': 'features'}, inplace=True)

# Visualize the DataFrame
coeffs_df

```

```

[17]:

```

	features	mean_coeffs_sign	coeffs_std
12	quantitative_preproc__GrLivArea	21714.189242	5388.545828
2	quantitative_preproc__OverallQual	18998.323970	2305.063393
21	quantitative_preproc__GarageCars	4931.342986	3460.563500
5	quantitative_preproc__BsmtFinSF1	4832.247887	3768.298030
33	quantitative_median_preproc__MasVnrArea	3852.195967	1944.982645
..	...	...	...
11	quantitative_preproc__LowQualFinSF	-164.544205	406.404862

28	quantitative_preproc__PoolArea	-207.523591	2912.103439
17	quantitative_preproc__BedroomAbvGr	-221.778399	723.354349
18	quantitative_preproc__KitchenAbvGr	-338.259106	568.481478
0	quantitative_preproc__MSSubClass	-4577.270253	1537.700679

[222 rows x 3 columns]

```
[18]: # filter df to not show the coeffs that went to 0
filtered_coeffs = coeffs_df.loc[(coeffs_df['mean_coeffs_sign'] != 0.0) &
                                (coeffs_df['mean_coeffs_sign'] != -0.0)]

# Create a function of the graph to reuse later
def barplot(figsize, title, data, x, y, xlabel, ylabel, error_data=None):
    # Plot the graph
    plt.figure(figsize=figsize)


    # Title
    plt.title(title, fontsize=14)

    # Graph
    sns.barplot(data=data, x=x, y=y, ci=None, palette="Blues_d")

    # Adding error bars if error_data is provided
    if error_data is not None:
        std_err = error_data
        plt.errorbar(data[x], data[y], yerr=std_err, fmt='none', color='k',
                    ↪ capsize=5)

    # Label
    plt.xticks(rotation=90)
    plt.xlabel(xlabel)
    plt.ylabel(ylabel)
    plt.show()

# Example usage:
barplot((18, 8), 'Coefficients direction (sign)', filtered_coeffs,
        'features', 'mean_coeffs_sign', 'Features', 'Beta (coeffs)',
        error_data=filtered_coeffs['coeffs_std'])
```




- Beta (coefficients) represents the gradient (slope) of the regression
- Beta is the change in y, per unit change in x (if all the other values stay the same)
- The positive values, if x increases, so does y
- The negative values, if x increases, y decreases

The Beta related to the feature “GrLivArea” > Beta related to the feature “OverallQual”, in other words “GrLivArea” has a greater contribution than “OverallQual” to the target value.

Let’s compare the absolute value of each coefficient.

```
[19]: # extract the absolute value of each coeff
filtered_coeffs['mean_coeffs_sign'] = filtered_coeffs['mean_coeffs_sign'].abs()
filtered_coeffs.sort_values(by=['mean_coeffs_sign'], ascending=False,
                             inplace=True)

barplot((18, 8), 'Coefficients direction (sign)', filtered_coeffs,
        'features', 'mean_coeffs_sign',
        'Features', 'Beta (coeffs)',
        error_data=filtered_coeffs['coeffs_std'])
```



This is what we normally use as a value of feature importance (feature selection methods, Lasso model automatically does this for us). Initially with more than 200 features, we ended with these, which are the most important for the model.

### 1.7.1 Calculate the t-statistic and p-value for each coefficient

The t-value, also known as the t-test, is an important focus of attention, since it is the link that tells us whether the association between an explanatory variable and the response is statistically significant. The t-value is simply the estimate/standard error, and thus can be interpreted as the distance of the estimate from 0, measured in standard errors. Given a t-value and sample size, the software can provide an accurate p-value; for large samples, t-values greater than 2 or less than -2 correspond to  $p < 0.05$ , although these thresholds are higher for smaller sample sizes.

- t-test tests the null hypothesis:  $b=0$
- Tests how big  $b$  is, compared to its variability

If  $t$  is too big or too small  $\rightarrow$  the probability that  $b=0$  is small, then, the regression coefficient is statistically significant.

```
[20]: filtered_coeffs['t_test'] = filtered_coeffs['mean_coeffs_sign'] / \
      filtered_coeffs['coeffs_std']
      filtered_coeffs
```

```

[20]:
      features  mean_coefs_sign \
12      quantitative_preproc__GrLivArea      21714.189242
 2      quantitative_preproc__OverallQual      18998.323970
21      quantitative_preproc__GarageCars      4931.342986
 5      quantitative_preproc__BsmtFinSF1      4832.247887
 0      quantitative_preproc__MSSubClass      4577.270253
33      quantitative_median_preproc__MasVnrArea      3852.195967
20      quantitative_preproc__Fireplaces      2968.611112
 8      quantitative_preproc__TotalBsmtSF      2396.445050
22      quantitative_preproc__GarageArea      2317.755954
 1      quantitative_preproc__LotArea      2305.309418
209      ordinal_quality_preproc__HeatingQC      2166.741426
208      ordinal_quality_preproc__BsmtCond      2132.256908
205      ordinal_quality_preproc__ExterQual      2105.917008
207      ordinal_quality_preproc__BsmtQual      2061.159714
210      ordinal_quality_preproc__KitchenQual      1832.783503
 4      quantitative_preproc__YearBuilt      1736.087225
19      quantitative_preproc__TotRmsAbvGrd      1702.281167
212      ordinal_quality_preproc__GarageCond      1694.193995
 9      quantitative_preproc__1stFlrSF      1653.541318
203      qualitative_mode_preproc__GarageType_Detchd      1596.794275
204      qualitative_mode_preproc__GarageType_nan      1531.679661
202      qualitative_mode_preproc__GarageType_CarPort      1472.624617
23      quantitative_preproc__WoodDeckSF      1367.713771
211      ordinal_quality_preproc__GarageQual      1324.697914
206      ordinal_quality_preproc__ExterCond      1301.277717
213      ordinal_exposure_preproc__BsmtExposure      1034.914082
214      ordinal_finish_preproc__BsmtFinType1      921.930767
35      quantitative_mode_preproc__YearRemodAdd      832.259605
200      qualitative_mode_preproc__GarageType_Basment      802.636039
13      quantitative_preproc__BsmtFullBath      691.008010
201      qualitative_mode_preproc__GarageType_BuiltIn      648.813157
27      quantitative_preproc__ScreenPorch      572.670701
216      ordinal_garage_preproc__GarageFinish      563.776395
215      ordinal_finish_preproc__BsmtFinType2      560.662227
15      quantitative_preproc__FullBath      522.115608
 3      quantitative_preproc__OverallCond      510.736195
34      quantitative_mode_preproc__LotFrontage      469.228607
36      quantitative_mode_preproc__GarageYrBlt      447.640997
217      qualitative_preproc__Condition2_RRAe      424.845662
18      quantitative_preproc__KitchenAbvGr      338.259106
17      quantitative_preproc__BedroomAbvGr      221.778399
28      quantitative_preproc__PoolArea      207.523591
198      qualitative_mode_preproc__GarageType_2Types      196.123903
218      qualitative_preproc__RoofMatl_Metal      189.083249
11      quantitative_preproc__LowQualFinSF      164.544205
32      quantitative_median_preproc__LotFrontage      132.491923


```

196	qualitative_mode_preproc__Electrical_SBrkr	88.227701
16	quantitative_preproc__HalfBath	50.949524
25	quantitative_preproc__EnclosedPorch	49.487195
219	qualitative_preproc__Exterior2nd_AsphShn	33.511129
24	quantitative_preproc__OpenPorchSF	25.551062
6	quantitative_preproc__BsmtFinSF2	9.872694
26	quantitative_preproc__3SsnPorch	7.166088
220	qualitative_preproc__Exterior2nd_Other	4.998262
30	quantitative_preproc__MoSold	3.972144
14	quantitative_preproc__BsmtHalfBath	2.744620

	coeffs_std	t_test
12	5388.545828	4.029694
2	2305.063393	8.241996
21	3460.563500	1.425012
5	3768.298030	1.282342
0	1537.700679	2.976698
33	1944.982645	1.980581
20	1423.609182	2.085271
8	2727.333788	0.878677
22	2483.236836	0.933361
1	1463.212331	1.575513
209	3144.635952	0.689028
208	3213.690150	0.663492
205	3196.704900	0.658777
207	3127.151470	0.659117
210	2825.582280	0.648639
4	1338.119735	1.297408
19	2364.765482	0.719852
212	2746.198706	0.616923
9	1642.856198	1.006504
203	3009.293957	0.530621
204	2687.776318	0.569869
202	3011.198185	0.489049
23	1080.613787	1.265682
211	2194.282506	0.603704
206	2586.492725	0.503105
213	2001.205826	0.517145
214	2247.865856	0.410136
35	990.156692	0.840533
200	2347.665914	0.341887
13	914.937378	0.755252
201	2241.082752	0.289509
27	684.686922	0.836398
216	1645.095017	0.342701
215	1571.393400	0.356793
15	997.536761	0.523405

3	624.005094	0.818481
34	902.797409	0.519750
36	640.038983	0.699396
217	1396.798415	0.304157
18	568.481478	0.595022
17	723.354349	0.306597
28	2912.103439	0.071262
198	1235.710207	0.158714
218	736.645253	0.256682
11	406.404862	0.404878
32	504.524588	0.262607
196	877.854544	0.100504
16	245.818373	0.207265
25	203.050524	0.243719
219	252.710121	0.132607
24	177.900767	0.143625
6	69.226082	0.142615
26	42.233511	0.169678
220	44.351484	0.112697
30	103.620553	0.038334
14	27.308623	0.100504

```
[21]: barplot((18, 8), 't-test Plot', filtered_coeffs,
            'features', 't_test',
            'Features', 't-test')
```



Based on the t values, we can, for example, trust the value of the coefficient of variable “OverallQual” more than that of variable “GrLivArea”.

## 1.7.2 Effect plots

Features whose coefficient is bigger are said to be more important. These coefficients alone don't determine the value of the target, to have this, we need the combination of these coefficients with the value of the variable, the one that contributes to the final target outcome.

So, instead of plot just the coefficients we can plot the effects.

```
[22]: # get the coefficients and the values of the variables
scaler_names = grid_search.best_estimator_.named_steps.preprocessor.
    ↪get_feature_names_out()
final_coeffs = grid_search.best_estimator_.named_steps.regressor.coef_
X_train_transformed = grid_search.best_estimator_.named_steps['preprocessor'].
    ↪transform(X_train)

# calculate the effects
effects = final_coeffs * X_train_transformed

# dividing the variables in groups of 50
num_variables = len(scaler_names)
```



```

group_size = 50
num_groups = int(np.ceil(num_variables / group_size))

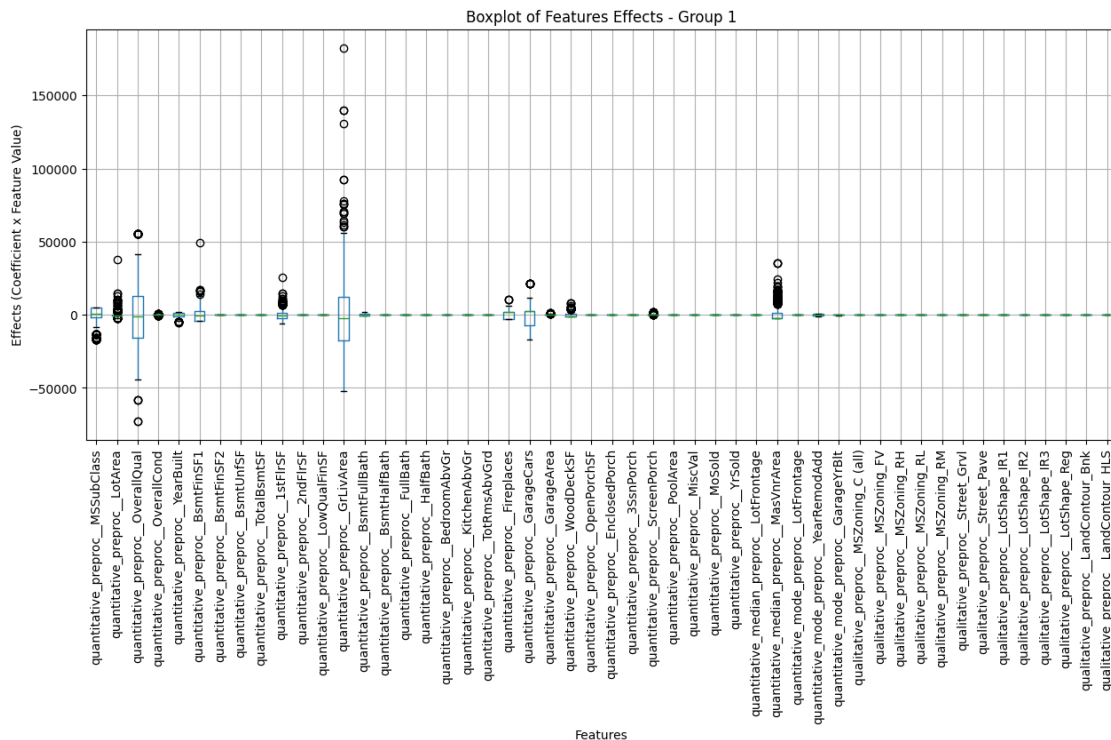
# iterate for the variables groups
for i in range(num_groups):
    start_index = i * group_size
    end_index = min((i + 1) * group_size, num_variables)

    # select the variables for the actual group
    scaler_names_group = scaler_names[start_index:end_index]
    effects_group = effects[:, start_index:end_index]


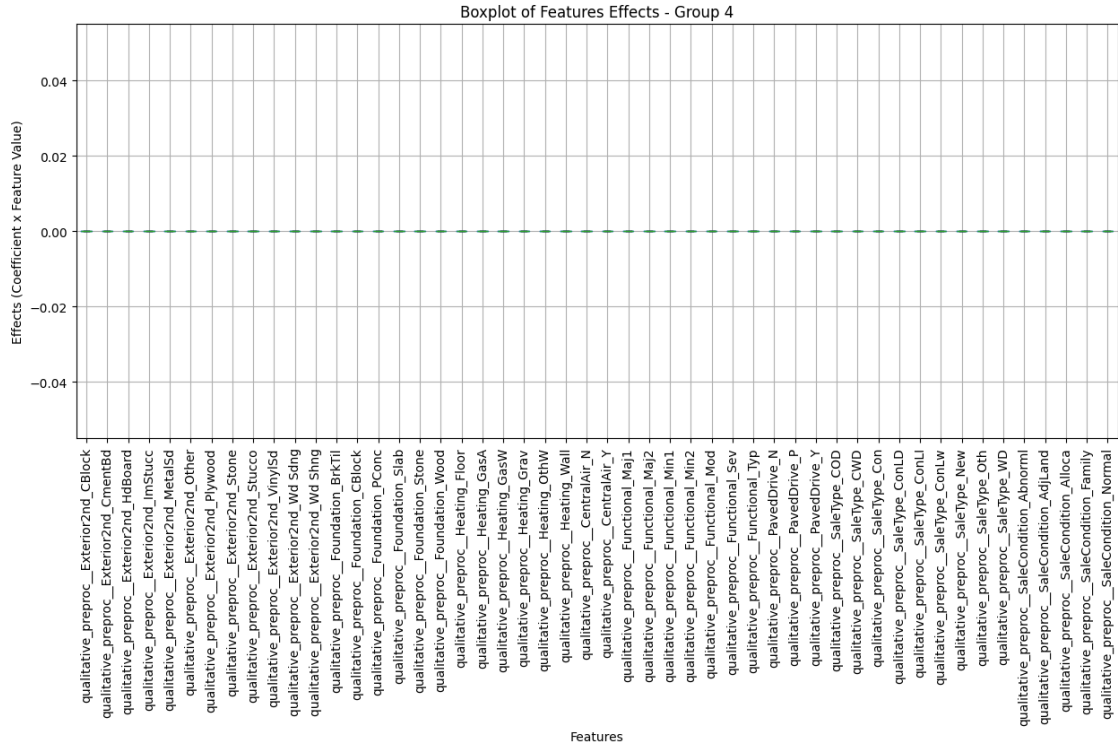
    # create the dataframe for the actual group
    effects_df_group = pd.DataFrame(effects_group, columns=scaler_names_group)

    # plot the boxplot for the actual group
    plt.figure(figsize=(15, 6))
    effects_df_group.boxplot()
    plt.title(f'Boxplot of Features Effects - Group {i+1}')
    plt.ylabel('Effects (Coefficient x Feature Value)')
    plt.xlabel('Features')
    plt.xticks(rotation=90)
    plt.show()

```







As we can see, the main drivers of the house price are the variables “OverallQual”, “GrLivArea”, “ExterQual”, “KitchenQual”, “BsmtExposure”.


## 1.8 Local interpretability

The same interpretation that we did globally (for all instances) we can also do locally for just one instance if you are interested in a specific case.

```
[24]: # calculate the effects of instance 1000
effects1000 = final_coefs * X_train_transformed[1000]

# create the dataframe for the instance 1000
effects_df1000 = pd.DataFrame({'mean_coefs_sign': final_coefs},
                               index=scaler_names).reset_index()
effects_df1000.sort_values(by=['mean_coefs_sign'], ascending=False,
                           inplace=True)
effects_df1000.rename(columns={'index': 'features'}, inplace=True)
effects_df1000 = effects_df1000.loc[
    (effects_df1000['mean_coefs_sign'] != 0.0) &
    (effects_df1000['mean_coefs_sign'] != -0.0)]

# plot the bar graph
barplot((18, 8), 'Effects Plot for Variable 1000', effects_df1000,
        'features', 'mean_coefs_sign',
        'Features', 'Effects (Coefficient x Feature Value)')
```



As we can see, these are the effects of the most important variables on the final price of the index house equal to 1000.

## 1.9 Conclusion

Here I want to summarize the main advantages of the linear regression models and also some of the limitations.

Advantages:

- Predict the target as a linear combination of the predictors (weighted sum of the predictors) and as humans we are very good interpreting linear models.
- We can use statistical tests to decide if / how much we can trust the model and its parameters ( i)
- Intrinsically explainable by design
- We can use regularization to reduce the feature space
  - Optimize for interpretability

Limitations:

- The interpretation of the weight / coefficient is contrastive (depends on all other features)
  - Features with positive correlation coefficient show a negative weight
- Make assumptions on the data -> when they are not met, we can't trust the model
- Multicollinearity affects interpretability
- Interactions between the variables won't be captured

Credits:

<https://www.trainindata.com/>